

1. Operating Systems Tutorial

Formalities, OS and C Basics, ABIs, Linking

Péter Bohner | 2023-11-02

OKAY, HUMAN.

HUH?

BEFORE YOU
HIT 'COMPILE',
LISTEN UP.

YOU KNOW WHEN YOU'RE
FALLING ASLEEP, AND
YOU IMAGINE YOURSELF
WALKING OR
SOMETHING,

AND SUDDENLY YOU
MISSTEP, STUMBLE,
AND JOLT AWAKE?

YEAH!

WELL, THAT'S WHAT A
SEGFALT FEELS LIKE.

DOUBLE-CHECK YOUR
DAMN POINTERS, OKAY?

Inhaltsverzeichnis

1. Organisation
2. OS Basics
3. C Basics
4. Using the template
5. ABIs
6. Linking
7. End

Organisation
oooo

OS Basics
oooooo

C Basics
oooooooooooo

Using the template
o

ABIs
oo

Linking
ooo

End
o

Introduction

About Me

- Péter Böhner (`peter.bohner@student.kit.edu`)
- CS Bachelor, 5. Semester

Introduce yourselves

- Name?
- Field of Study?
- Am I your regular tutor?
- C/Systems programming/operating systems knowledge?
- Experience with Linux?
- Expectations for the tutorial?

Links

- Everything is in ILIAS
- My website https://bohner.me/teaching/ws2023_ostut/
- ATIS: <https://www.atis.informatik.kit.edu/513.php>
- Compiler Explorer: <https://godbolt.org/>
- Feedback: <https://forms.gle/9CwJSKidKibubran9>

Where should I ask my questions?

- 1 Ask in the tutorial
- 2 ILIAS-Forum
- 3 Tutor (email: peter.bohner@student.kit.edu, discord: xzvf)
- 4 Mail to Peter Maucher
- 5 Mail to prof

- Exam: Three exercises (20P each) with mixed theory and practice (programming) Exact modalities unknown **different from past exams!**
- No separate exam for repeat takers, probably
- Exercises optional, but highly recommended
 - Complete short introductory tasks!
 - Tackle larger tasks if interested
- The contents of the exercises are part of the exam, not just the lecture!

Completing the exercises

- Submissions must compile and run on ****Linux x86_64****
- Exercises with x86 assembly and linux syscalls

your host	recommended development environment
* x86_32	ATIS
Linux x86_64	your host :-)
*BSD x86_64	Linux VM/dual boot
Windows x86_64	WSL2 or Linux VM/dual boot
Windows ARM	ATIS
MacOS ARM	ATIS, Asahi Linux
MacOS x64	Linux VM/dual boot
Sonstiges	ATIS

Tasks and responsibilities of operating systems

Abstraction/Standardization

- Devices and how to talk to them differ greatly
 - Remember the Driver-CDs“ shipped with motherboards back in the days so you could properly configure and run things?
- Strange user demands: Programs should run on more than one hardware configuration
 - Abstract away hardware details!

Tasks and responsibilities of operating systems

Resource management

- You want to print? Too bad, another program is already using that printer.
- You want to access the storage drive? Too bad, another program is already doing that.
- You want to get CPU time? Too bad, this while (true) loop is more important.

Tasks and responsibilities of operating systems

Security and Protection

- You are a good citizen and use a password manager. What could happen without your helpful OS? Other programs may read its memory!
- You copied a password to your clipboard? Oops. You're on your own there. The Clipboard is not provided by the OS and mostly has no special protections.
- You write a cool little program that fills a buffer with a random value. Sadly you made a mistake and missed a bounds check. What happens? You crash, but your text editor doesn't suddenly have its memory overwritten!

Tasks and responsibilities of operating systems

Provide an execution environment for applications

- What does that mean? Basically all of the above combined and more. Make a homely place where applications like to live!

Organisation
oooo

OS Basics
ooo●oo

C Basics
oooooooooooo

Using the template
o


ABIs
oo

Linking
ooo

End
o

Kernel / User mode differences

- What are the differences between a processor running in kernel or user mode?
- Why are both modes needed?



- In kernel mode you have full access to privileged instructions.

Organisation
○○○○

OS Basics
○○○○●○

C Basics
○○○○○○○○○○○○

Using the template
○

ABIs
○○

Linking
○○○

End
○

What instructions are privileged?

- Change control registers for memory mappings
 - Read other process's memory
- Disable / Enable interrupts
 - No preemption for you.
- Access platform devices (network card, storage, printer, ...)
- Some nice registers: LGDT (Load Global Descriptor Table) or the LLDT, INVD (Invalidate cache), HLT (Halt processor!)
- Is **MOV** (Move) a privileged instruction?
 - Yes, if moving to debug/control registers or privileged memory locations!

The C Programming Language

History

- Developed by Dennis Ritchie at Bell Labs in 1972
- successor to B
- one of the most used (systems-) programming languages

Properties

- imperative, procedural, not OO
- low level
- manual memory management
 - fun with pointers
 - CVE-Factory

Datatypes - Primitives

name	minimum size (bytes)	x86_64 Linux
char	1	1
short	2	2
int	2	4
long	4	8
long long	8	8
float	-	4
double	-	8
long double	-	16

- Modifier: unsigned signed, all types are signed by default
- except char, that is architecture dependent . . .
- except char which is implementation dependant

Datatypes - fixed size

```
#include<inttypes.h> // or #include<stdint.h>
int8_t a; uint16_t b;
int_least16_t c; int_fast16_t d;
intptr_t p; ssize_t ss; size_t s;
```

Datatypes - boolean

- does not really exist (C99 `_Bool` type)
- use 0 / 1 (0 = false, nonzero true)
- or `stdbool.h`

```
// #include<stdbool.h>
typedef unsigned char bool;
#define false 0
#define true 1
int main () {
    bool hey = true;
    return 0;
}
```

- Size of types/variables

```
int foo; printf("%d", sizeof(foo));
```


Datatypes - Arrays

```
int arr[5] = {1, 2, 3, 4, 5};  
int arr2[] = {1, 2, 3, 4, 5};  
int arr4[10] = {1,2};  
int arr3[5];
```

- continuous memory area
- size is part of type, `sizeof(int[5])` works correctly
- How do you get the size of an array during runtime?
- That's the neat part: you don't :D
 - `sizeof(arr)` is `sizeof(int *)`
 - because arrays are just syntax sugar for pointers

Datatypes - Strings

```
char *string = "mystring";  
char[] s2 = "another";  
char[] s3 = {'H', 'i', '!', '\\0'};
```

- strings are char arrays with NULL terminator at the end
- One byte longer than number of ASCII characters
- Uses ascii codes

Datatypes - Structs

```
struct Person {
    int age;
    char *name;
};

typedef struct Person person_t;

int main(void) {
    struct Person peter = {.name = "Peter", .age = 20};
    person_t also_peter = peter;
    printf("%s is %d years old.\n", peter.name, peter.age);
}
```

- Product type
- Not a class: No hidden vtables/inheritance, no methods.

Locals and globals

- What is the difference between local and global variables?
- **globals**: Live while the program runs. Placed in data segments.
- **locals**: Are only accessible in scope of declaration. Live on stack or in registers

```
int uninitialized_global; // 0
int initialized_global = 0xCAFE; // 1
const int const_global = 0xBEEF; // 2
int func() {int local; /**/}
```

Where is each symbol stored?

- 1 .bss
- 2 .data
- 3 .rodata
- 4 stack (or register)

Pointers

```
typedef struct {int age;char *name;} Person;
void foo(Person p) { p.age = 100; printf("%d ", p.age);}
void bar(Person *p) {p->age = 200; printf("%d ", p->age);}
int main(void) {
    Person p = {.name = "Me", .age = 20};
    printf("%d ", p.age);
    foo(p);
    printf("%d ", p.age);
    bar(&);
    printf("%d ", p.age);
}
```

output?

20 100 20 200 200

Pointers

- A variable that points to an arbitrary region of memory
- a `int *` points to an `int`, `void *` points to anything
- the element size of `void` is **undefined**
- Java references are like pointers, but without pointer arithmetic

Pointers

```
void swap(int *a, int *b) {
    int tmp = *a; //dereference with *
    *a = *b; *b = tmp;
}

int main(void) {
    int x = 5, y = 10;
    printf("x = %d; y = %d\n", x, y);
    swap(&a, &b); //take address with &
    printf("x = %d; y = %d\n", x, y);
}
```

output?

```
x = 5; y = 10
x = 10; y = 5
```

Organisation
○○○○

OS Basics
○○○○○○

C Basics
○○○○○○○○○○●○

Using the template
○

ABIs
○○

Linking
○○○

End
○

Pointers

```
typedef struct {char *name; int age; } Person;
void foo(Person *p) {
    printf("%s is %d years old\n", p->name; (*p).age); //-> is just syntax sugar
}
void bar(Person people[], size_t count) {
    for(size_t i = 0; i < count; i++) {
        printf("%s is %d years old\n", people[i].name, (people + i)->age); //-> as is []
    }
}
```

Operators applicable on int*

* & -> [] + -

Demo time

Organisation
oooo

OS Basics
oooooo

C Basics
oooooooooooooooo

Using the template
●

ABIs
oo

Linking
ooo

End
o

What is an ABI?

Application Binary Interface

What does that specify?

- Interface of **binary** programs (i.e. after compilation)
- Instruction Set (e.g. x68, ARM)
- Calling convention (e.g. **cdecl** or **System V AMD64 ABI**)
- Basic data types and their size / alignment (int, sizeof(int))
- How to perform System Calls

ABI of a library

This tem is used when talking about binary compatibility of different library versions. -> Do you need to recompile your code against the new version?

Calling conventions

What's the usual calling convention on modern Linux?

System V AMD64 ABI

- Integer arguments in `rdi`, `rsi`, `rdx`, `rcx`, `r8`, `r9`, then stack
- FP arguments in `xmm0` to `xmm7`
- Integer return value in `rax`, `rdx`

<https://godbolt.org/z/68xexn>

Static vs Dynamic linking

What is the difference between static and dynamic linking?

- A static library is a collection of object files
 - The linker can treat it as normal code
- A dynamic library is loaded and linked at runtime

Pros and cons

- 1 static Linking
 - + Unused references can be elided
 - + Library calls just as fast as local ones
 - + No runtime overhead for loading and relocation
 - - Library can not be shared -> Memory overhead
- 2 dynamic linking
 - + Library (Code segment) can be shared

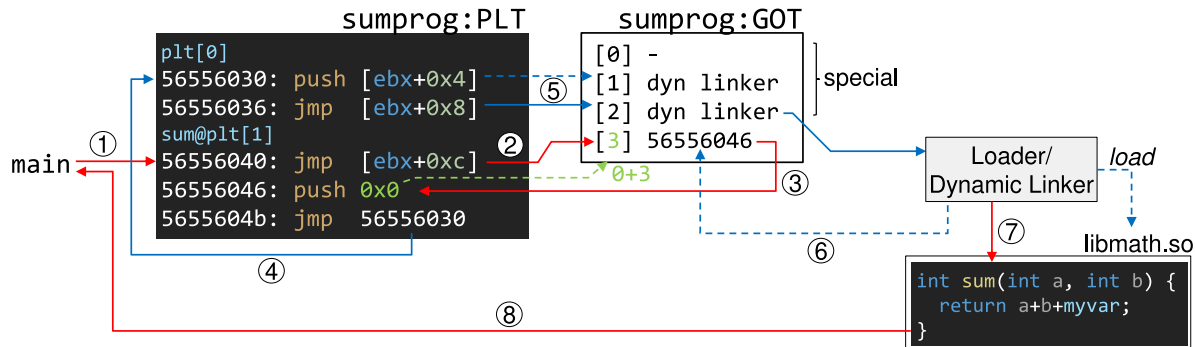
What does PIC stand for?

Position Independent Code

Independent to what?

It's position in the process' address space.

- Shared libraries are loaded somewhere in the address space of a process
- PIC uses relative addresses only -- and can therefore be loaded anywhere
- How do you find global symbols then?
 - GOT, the Global Offset Table



End

Are there any questions?

Leave anonymous feedback



<https://forms.gle/9CwJSKidKibubran9>

Bye!

Until next time, **Wed, 8.11.2023!**

Organisation
○○○○

OS Basics
○○○○○

C Basics
○○○○○○○○○○○○

Using the template
○

ABIs
○○

Linking
○○○

End
●