

## Scheduling basics

---

**What is a Scheduler? Why do we even need it?**

## What is a Scheduler? Why do we even need it?

- Maps processes to resources

## What is a Scheduler? Why do we even need it?

- Maps processes to resources
- Ideally: Every process gets what it needs some day

## What is a Scheduler? Why do we even need it?

- Maps processes to resources
- Ideally: Every process gets what it needs some day

## What Schedulers do you know?

## What is a Scheduler? Why do we even need it?

- Maps processes to resources
- Ideally: Every process gets what it needs some day

## What Schedulers do you know?

- CPU-Scheduler: The classic

## What is a Scheduler? Why do we even need it?

- Maps processes to resources
- Ideally: Every process gets what it needs some day

## What Schedulers do you know?

- CPU-Scheduler: The classic
- Disk-Scheduler: Why have one?

## What is a Scheduler? Why do we even need it?

- Maps processes to resources
- Ideally: Every process gets what it needs some day

## What Schedulers do you know?

- CPU-Scheduler: The classic
- Disk-Scheduler: Why have one? Multiplexing but also efficiency!



## What is a Scheduler? Why do we even need it?

- Maps processes to resources
- Ideally: Every process gets what it needs some day

## What Schedulers do you know?

- CPU-Scheduler: The classic
- Disk-Scheduler: Why have one? Multiplexing but also efficiency!
- Network I/O:

## What is a Scheduler? Why do we even need it?

- Maps processes to resources
- Ideally: Every process gets what it needs some day

## What Schedulers do you know?

- CPU-Scheduler: The classic
- Disk-Scheduler: Why have one? Multiplexing but also efficiency!
- Network I/O: When to send packets, which packets to drop, QoL,...

**What are the differences? When are they used?**

### What are the differences? When are they used?

- LTS: Decide which processes to put in the *run queue*

### What are the differences? When are they used?

- LTS: Decide which processes to put in the *run queue*
- STS: Decide which process runs on the *CPU*

## What are the differences? When are they used?

- LTS: Decide which processes to put in the *run queue*
  - STS: Decide which process runs on the *CPU*
  - MTS: Temporarily removes processes from main memory (and e.g. writes them out to disk)
- ⇒ Reduce degree of multiprogramming, make room in memory (and a few other reasons)







# Process states



# Process states

new

ready

running

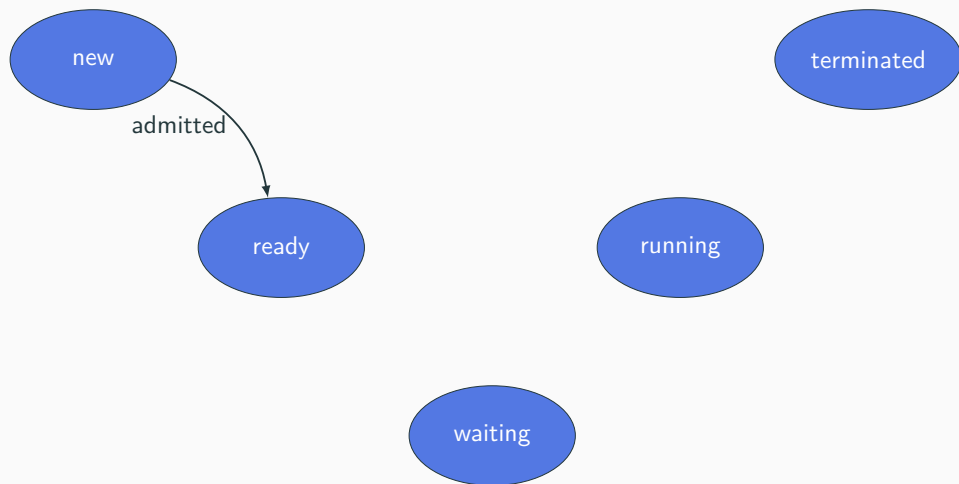
# Process states




# Process states




# Process states




## Process states



# Process states

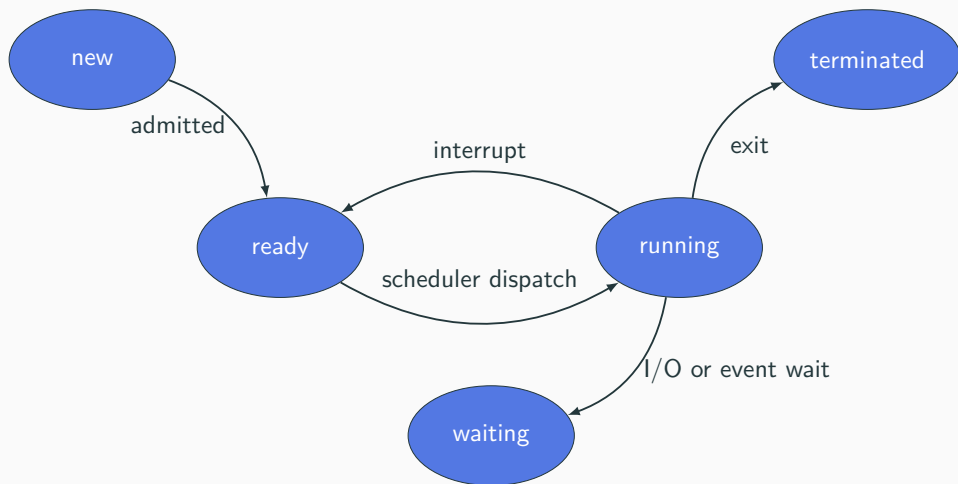


# Process states





## Process states



„I/O or event wait“? When does a process move from ready to waiting?

- Network / Disk I/O

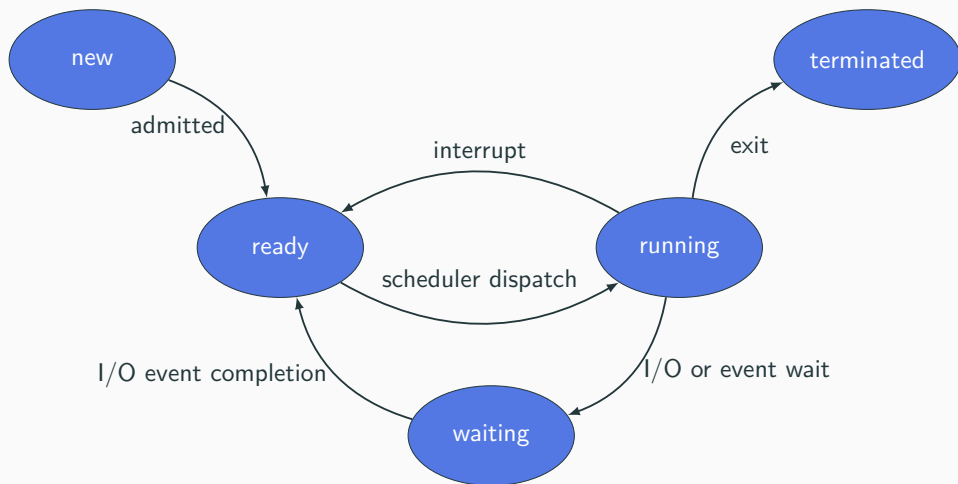
**„I/O or event wait“? When does a process move from ready to waiting?**

- Network / Disk I/O
- Mutex or other inter-process synchronisation

**„I/O or event wait“? When does a process move from ready to waiting?**

- Network / Disk I/O
- Mutex or other inter-process synchronisation
- Sleepyness

# Process states



**What makes a good Scheduler good?**

Let's play scheduler!

## What makes a good Scheduler good?

Let's play scheduler!

## Some metrics

- Processor utilization: Percentage of working time
- Throughput: How many jobs do you finish?
- Turnaround time: Wallclock-time from submission to finish
- Waiting time: How long did it spend in the ready queue
- Response time: Time between submission of a request and first response (e.g. key press to echo on screen)

**What does your hardware need to support to allow non-cooperative scheduling?**



**What does your hardware need to support to allow non-cooperative scheduling?**

Timer Interrupts! Waiting for a cosmic ray to hit, a network package to arrive, a system call or any other random interrupt gets old fast :)

## Scheduling - When to interrupt

Any guesses for how long a timeslice usually is?

### Any guesses for how long a timeslice usually is?

2ms - 200ms

- On windows it depends on the configuration (favor foreground / background processes)

### Any guesses for how long a timeslice usually is?

2ms - 200ms

- On windows it depends on the configuration (favor foreground / background processes)

Which setting has the longer timeslice?

### Any guesses for how long a timeslice usually is?

2ms - 200ms

- On windows it depends on the configuration (favor foreground / background processes)

Which setting has the longer timeslice?

20ms to 30ms for foreground, 180ms to 200ms for background

### Any guesses for how long a timeslice usually is?

2ms - 200ms

- On windows it depends on the configuration (favor foreground / background processes)

Which setting has the longer timeslice?

20ms to 30ms for foreground, 180ms to 200ms for background

- Linux's „Completely Fair Scheduler“ adjusts them dynamically based on the priority, number of processes, ...

## Any guesses for how long a timeslice usually is?

2ms - 200ms

- On windows it depends on the configuration (favor foreground / background processes)

Which setting has the longer timeslice?

20ms to 30ms for foreground, 180ms to 200ms for background

- Linux's „Completely Fair Scheduler“ adjusts them dynamically based on the priority, number of processes, ...

## Benefits of shorter/longer timeslices?

## Any guesses for how long a timeslice usually is?

2ms - 200ms

- On windows it depends on the configuration (favor foreground / background processes)

Which setting has the longer timeslice?

20ms to 30ms for foreground, 180ms to 200ms for background

- Linux's „Completely Fair Scheduler“ adjusts them dynamically based on the priority, number of processes, ...

## Benefits of shorter/longer timeslices?

- Short: High interactivity, higher overhead



## Any guesses for how long a timeslice usually is?

2ms - 200ms

- On windows it depends on the configuration (favor foreground / background processes)

Which setting has the longer timeslice?

20ms to 30ms for foreground, 180ms to 200ms for background

- Linux's „Completely Fair Scheduler“ adjusts them dynamically based on the priority, number of processes, ...

## Benefits of shorter/longer timeslices?

- Short: High interactivity, higher overhead
- Long: Lower interactivity, smaller overhead

Pitfalls - How would you implement this?

### Pitfalls - How would you implement this?

- You would need to have future knowledge to figure out the job length!

### Pitfalls - How would you implement this?

- You would need to have future knowledge to figure out the job length!  
How do you solve this?

## Pitfalls - How would you implement this?

- You would need to have future knowledge to figure out the job length!  
How do you solve this?
- **Predict** the *future* based on *past* behaviour

### Pitfalls - How would you implement this?

- You would need to have future knowledge to figure out the job length!  
How do you solve this?
- **Predict** the *future* based on *past* behaviour
- Does this work?

## Pitfalls - How would you implement this?

- You would need to have future knowledge to figure out the job length!  
How do you solve this?
- **Predict** the *future* based on *past* behaviour
- Does this work?  
„*THERE IS AS YET INSUFFICIENT DATA FOR A MEANINGFUL ANSWER.*“  
~ Isaac Asimov, „*The Last Question*“ (Comic)

## Pitfalls - How would you implement this?

- You would need to have future knowledge to figure out the job length!  
How do you solve this?
- **Predict** the *future* based on *past* behaviour
- Does this work?  
„*THERE IS AS YET INSUFFICIENT DATA FOR A MEANINGFUL ANSWER.*“  
~ Isaac Asimov, „*The Last Question*“ (Comic)  
You need some balanced initial value. Not *that* big of a deal with preemption though. Why?



## Pitfalls - How would you implement this?

- You would need to have future knowledge to figure out the job length!  
How do you solve this?

- **Predict** the *future* based on *past* behaviour

- Does this work?

„*THERE IS AS YET INSUFFICIENT DATA FOR A MEANINGFUL ANSWER.*“

~ Isaac Asimov, „**The Last Question**“ ([Comic](#))

You need some balanced initial value. Not *that* big of a deal with preemption though. Why?


Interrupt the process after the estimated time is over.

What is priority scheduling? Why would you use it?

## What is priority scheduling? Why would you use it?

- Each process is assigned a priority
- The process with the highest priority is chosen

# Multi-Level Feedback Queues



## How it works

- All processes start in the highest queue
  - When they use up their timeslice and are preempted, they descend
  - If they block before, they stay in the level (optionally: Are moved up)
- ⇒ I/O bound processes rise to the top and react quickly, CPU bound processes get longer timeslices but less often

# Multi-Level Feedback Queues

Short time slice

		C	D
--	--	---	---

A/B used their full timeslice

A	B		
---	---	--	--

Long time slice

--	--	--	--

## How it works

- All processes start in the highest queue
  - When they use up their timeslice and are preempted, they descend
  - If they block before, they stay in the level (optionally: Are moved up)
- ⇒ I/O bound processes rise to the top and react quickly, CPU bound processes get longer timeslices but less often

# Multi-Level Feedback Queues

Short time slice

		C	D
--	--	---	---

B didn't use it fully

	B		
--	---	--	--

A used its full timeslice

A			
---	--	--	--

## How it works

- All processes start in the highest queue
  - When they use up their timeslice and are preempted, they descend
  - If they block before, they stay in the level (optionally: Are moved up)
- ⇒ I/O bound processes rise to the top and react quickly, CPU bound processes get longer timeslices but less often

# A Flawless Scheduling Algorithm?

Consider the waiting time

# A Flawless Scheduling Algorithm?

## Consider the waiting time

- A few I/O-bound jobs could saturate the CPU!

⇒ The lower-level processes starve



# A Flawless Scheduling Algorithm?

## Consider the waiting time

- A few I/O-bound jobs could saturate the CPU!

⇒ The lower-level processes starve

## How could you fix this?

# A Flawless Scheduling Algorithm?

## Consider the waiting time

- A few I/O-bound jobs could saturate the CPU!

⇒ The lower-level processes starve

## How could you fix this?

- E.g. reset the whole thing after a given interval, so all start in the highest level again

# A Flawless Scheduling Algorithm?

## Consider the waiting time

- A few I/O-bound jobs could saturate the CPU!

⇒ The lower-level processes starve

## How could you fix this?

- E.g. reset the whole thing after a given interval, so all start in the highest level again
- "Boost" processes that waited for a long time

# A Flawless Scheduling Algorithm?

## Consider the waiting time

- A few I/O-bound jobs could saturate the CPU!

⇒ The lower-level processes starve

## How could you fix this?

- E.g. reset the whole thing after a given interval, so all start in the highest level again
- "Boost" processes that waited for a long time

## What metrics does it optimize?

- Utilization? Turnaround time? Throughput? Waiting time? Response time?

# A Flawless Scheduling Algorithm?

## Consider the waiting time

- A few I/O-bound jobs could saturate the CPU!

⇒ The lower-level processes starve

## How could you fix this?

- E.g. reset the whole thing after a given interval, so all start in the highest level again
- "Boost" processes that waited for a long time

## What metrics does it optimize?

- Utilization? Turnaround time? Throughput? Waiting time? Response time?
- Prefer I/O bound, prefer short jobs, group the rest based on their needs

## Process switching

---

### What does the kernel need to do when switching processes?

- Adjust Instruction, Base and Frame Pointer -- okay

### What does the kernel need to do when switching processes?

- Adjust Instruction, Base and Frame Pointer -- okay

The process shouldn't notice it was interrupted! What else do you need to adjust?



### What does the kernel need to do when switching processes?

- Adjust Instruction, Base and Frame Pointer -- okay

The process shouldn't notice it was interrupted! What else do you need to adjust?

- General purpose register

### What does the kernel need to do when switching processes?

- Adjust Instruction, Base and Frame Pointer -- okay

The process shouldn't notice it was interrupted! What else do you need to adjust?

- General purpose register
- Address space

### What does the kernel need to do when switching processes?

- Adjust Instruction, Base and Frame Pointer -- okay

The process shouldn't notice it was interrupted! What else do you need to adjust?

- General purpose register
- Address space
- And a bit of housekeeping:

### What does the kernel need to do when switching processes?

- Adjust Instruction, Base and Frame Pointer -- okay

The process shouldn't notice it was interrupted! What else do you need to adjust?

- General purpose register
- Address space
- And a bit of housekeeping: Open files, scheduling priorities, ...

⇒ These things make up the PCB - The

### What does the kernel need to do when switching processes?

- Adjust Instruction, Base and Frame Pointer -- okay

The process shouldn't notice it was interrupted! What else do you need to adjust?

- General purpose register
- Address space
- And a bit of housekeeping: Open files, scheduling priorities, ...

⇒ These things make up the PCB - The Process Control Block

### What does the kernel need to do when switching processes?

- Adjust Instruction, Base and Frame Pointer -- okay

The process shouldn't notice it was interrupted! What else do you need to adjust?

- General purpose register
- Address space
- And a bit of housekeeping: Open files, scheduling priorities, ...

⇒ These things make up the PCB - The Process Control Block

### Where are the PCBs stored? In user or kernel space?

### What does the kernel need to do when switching processes?

- Adjust Instruction, Base and Frame Pointer -- okay

The process shouldn't notice it was interrupted! What else do you need to adjust?

- General purpose register
- Address space
- And a bit of housekeeping: Open files, scheduling priorities, ...

⇒ These things make up the PCB - The Process Control Block

### Where are the PCBs stored? In user or kernel space?

- Kernel space! Users shouldn't be able to modify them

**Are the PCBs always valid?**



**Are the PCBs always valid?**

No! Some parts (registers, PC, etc) only when the process is not running. Why?

### **Are the PCBs always valid?**

No! Some parts (registers, PC, etc) only when the process is not running. Why?  
They are saved when it is switched out.


## In pictures

P1 — Living life

P2

# Process switching

## In pictures




# Process switching

## In pictures



# Process switching

## In pictures



## As text

1. Transition to the kernel (How?

## As text

1. Transition to the kernel (How? Interrupt, Exception, Syscall)



## As text

1. Transition to the kernel (How? Interrupt, Exception, Syscall)
2. Save the context of the process. What was that again?

## As text

1. Transition to the kernel (How? Interrupt, Exception, Syscall)
2. Save the context of the process. What was that again? Registers, Program counter, Stack pointer, etc.

## As text

1. Transition to the kernel (How? Interrupt, Exception, Syscall)
2. Save the context of the process. What was that again? Registers, Program counter, Stack pointer, etc.
3. Where do you save it to? A per-process kernel stack typically (as it is often moved there automatically) and then move it to the PCB.

## As text

1. Transition to the kernel (How? Interrupt, Exception, Syscall)
2. Save the context of the process. What was that again? Registers, Program counter, Stack pointer, etc.
3. Where do you save it to? A per-process kernel stack typically (as it is often moved there automatically) and then move it to the PCB.
4. Restore the context of the next process

## As text

1. Transition to the kernel (How? Interrupt, Exception, Syscall)
2. Save the context of the process. What was that again? Registers, Program counter, Stack pointer, etc.
3. Where do you save it to? A per-process kernel stack typically (as it is often moved there automatically) and then move it to the PCB.
4. Restore the context of the next process
5. Leave kernel mode and transfer control to the PC of the next process

# Threads



**What are processes, address spaces and threads? How do they relate to each other?**

What are processes, address spaces and threads? How do they relate to each other?

- A thread is an *entity of execution*, the personification of control flow



**What are processes, address spaces and threads? How do they relate to each other?**

- A thread is an *entity of execution*, the personification of control flow
- A thread lives in an address space, i.e. all the addresses that it can access and the data that is stored there

**What are processes, address spaces and threads? How do they relate to each other?**

- A thread is an *entity of execution*, the personification of control flow
- A thread lives in an address space, i.e. all the addresses that it can access and the data that is stored there
- Thread + Address Space = Process

## Spawn a few threads using pthreads!

Write a small program that creates five threads using the pthread library. Each thread should print its number (e.g., Hello, I am 4) and the main program should wait for each thread to exit.

# Thread models — One To One

## One To One

`new Thread()`

`new Thread()`

Process

A diagram illustrating the One To One thread model. It features a large light gray rectangular area. On the left side of this area, the text 'new Thread()' is written twice, stacked vertically. On the right side, there is a smaller, outlined rectangular box. Inside this box, the word 'Process' is written. This visualizes that each thread created in this model corresponds to a separate operating system process.


## Problems?

# Thread models — One To One

## One To One

`new Thread()`


`new Thread()`



## Problems?

# Thread models — One To One


## One To One



## Problems?


# Thread models — One To One

## One To One



# Thread models — One To One

## One To One



## Problems?



Problems and benefits of *One To One*?

### Problems and benefits of *One To One*?

- + Scales with core count

### Problems and benefits of *One To One*?

- + Scales with core count
- + Conceptually easy — the OS does the hard stuff

### Problems and benefits of *One To One*?

- + Scales with core count
- + Conceptually easy — the OS does the hard stuff
- + Blocking does not affect other threads

## Problems and benefits of *One To One*?

- + Scales with core count
- + Conceptually easy — the OS does the hard stuff
- + Blocking does not affect other threads
- + Can piggy-back onto the OS scheduler

## Problems and benefits of *One To One*?

- + Scales with core count
- + Conceptually easy — the OS does the hard stuff
- + Blocking does not affect other threads
- + Can piggy-back onto the OS scheduler
- **Must** piggy-back onto the OS scheduler

## Problems and benefits of *One To One*?

- + Scales with core count
- + Conceptually easy — the OS does the hard stuff
- + Blocking does not affect other threads
- + Can piggy-back onto the OS scheduler
- **Must** piggy-back onto the OS scheduler
- Relatively high overhead due to context switches

## Problems and benefits of *One To One*?

- + Scales with core count
- + Conceptually easy — the OS does the hard stuff
- + Blocking does not affect other threads
- + Can piggy-back onto the OS scheduler
- **Must** piggy-back onto the OS scheduler
- Relatively high overhead due to context switches
- Relatively high overhead *when creating one*



## Problems and benefits of *One To One*?

- + Scales with core count
- + Conceptually easy — the OS does the hard stuff
- + Blocking does not affect other threads
- + Can piggy-back onto the OS scheduler
- **Must** piggy-back onto the OS scheduler
- Relatively high overhead due to context switches
- Relatively high overhead *when creating one!!*

## Many to One

`new Thread()`

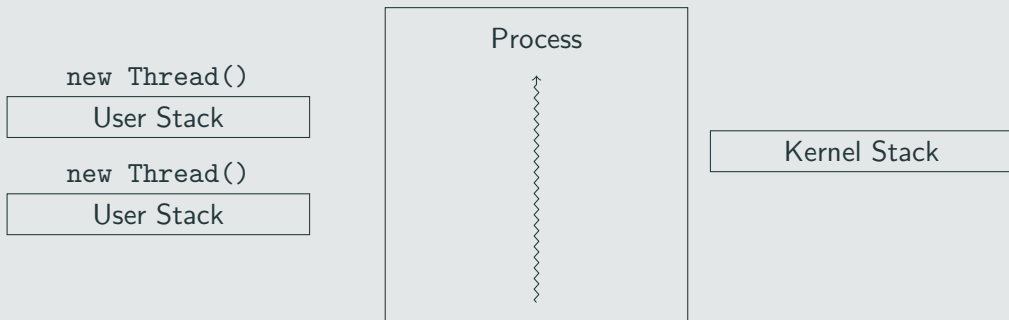
User Stack

Process




Kernel Stack


## Many to One




## Many to One




## Many to One



## Many to One



## Many to One



### Problems and benefits of *Many To One*?

Do they improve anything?



### Problems and benefits of *Many To One*?

Do they improve anything?

+ **Scales with core count?**

### Problems and benefits of *Many To One*?

Do they improve anything?

- Can only use one core

### Problems and benefits of *Many To One*?

Do they improve anything?

- Can only use one core

- + **Conceptually easy** — the OS does the hard stuff?

### Problems and benefits of *Many To One*?

Do they improve anything?

- Can only use one core
- Harder to implement — the OS doesn't help you much

### Problems and benefits of *Many To One*?

Do they improve anything?

- Can only use one core
- Harder to implement — the OS doesn't help you much
- + **Blocking does not affect other threads?**

### Problems and benefits of *Many To One*?

Do they improve anything?

- Can only use one core
- Harder to implement — the OS doesn't help you much
- Blocking *does* affect other threads

### Problems and benefits of *Many To One*?

Do they improve anything?

- Can only use one core
- Harder to implement — the OS doesn't help you much
- Blocking *does* affect other threads
- + **Can piggy-back onto the OS scheduler?**

### Problems and benefits of *Many To One*?

Do they improve anything?

- Can only use one core
- Harder to implement — the OS doesn't help you much
- Blocking *does* affect other threads
- Can *not* piggy-back onto the OS scheduler



## Problems and benefits of *Many To One*?

Do they improve anything?

- Can only use one core
- Harder to implement — the OS doesn't help you much
- Blocking *does* affect other threads
- Can *not* piggy-back onto the OS scheduler
- **Must** piggy-back onto the OS scheduler?

## Problems and benefits of *Many To One*?

Do they improve anything?

- Can only use one core
- Harder to implement — the OS doesn't help you much
- Blocking *does* affect other threads
- Can *not* piggy-back onto the OS scheduler
- + Can implement its **own** scheduler

## Problems and benefits of *Many To One*?

Do they improve anything?

- Can only use one core
- Harder to implement — the OS doesn't help you much
- Blocking *does* affect other threads
- Can *not* piggy-back onto the OS scheduler
- + Can implement its **own** scheduler
- **Relatively high overhead due to context switches?**

## Problems and benefits of *Many To One*?

Do they improve anything?

- Can only use one core
- Harder to implement — the OS doesn't help you much
- Blocking *does* affect other threads
- Can *not* piggy-back onto the OS scheduler
- + Can implement its **own** scheduler
- + Low overhead during context switches

## Problems and benefits of *Many To One*?

Do they improve anything?

- Can only use one core
- Harder to implement — the OS doesn't help you much
- Blocking *does* affect other threads
- Can *not* piggy-back onto the OS scheduler
- + Can implement its **own** scheduler
- + Low overhead during context switches
- **Relatively high overhead *when creating one?***

## Problems and benefits of *Many To One*?

Do they improve anything?

- Can only use one core
- Harder to implement — the OS doesn't help you much
- Blocking *does* affect other threads
- Can *not* piggy-back onto the OS scheduler
- + Can implement its **own** scheduler
- + Low overhead during context switches
- + Low overhead when creating one

# Thread models - Many To One

**Do you know a programming language / runtime using that?**

E.g. nodejs using its „event loop“

## **A small excursion - Structured Programming**

Control flow should fall into one of four patterns:

- Sequence: One block is executed after another

# Thread models - Many To One

**Do you know a programming language / runtime using that?**

E.g. nodejs using its „event loop“

## **A small excursion - Structured Programming**

Control flow should fall into one of four patterns:

- Sequence: One block is executed after another
- Selection: One or more are executed (i.e. an `if`-statement)



# Thread models - Many To One

**Do you know a programming language / runtime using that?**

E.g. nodejs using its „event loop“

## **A small excursion - Structured Programming**

Control flow should fall into one of four patterns:

- Sequence: One block is executed after another
- Selection: One or more are executed (i.e. an `if`-statement)
- Iteration: A block is executed more than once (i.e. a loop)

# Thread models - Many To One

**Do you know a programming language / runtime using that?**

E.g. nodejs using its „event loop“

## **A small excursion - Structured Programming**

Control flow should fall into one of four patterns:

- Sequence: One block is executed after another
- Selection: One or more are executed (i.e. an `if`-statement)
- Iteration: A block is executed more than once (i.e. a loop)
- Recursion: A block calls itself until an exit condition is met (i.e. recursion!)

Do you know any keyword in C which *doesn't* quite adhere to that but can instead totally spaghetti your control flow?

# Thread models - Many To One

Do you know a programming language / runtime using that?

E.g. nodejs using its „event loop“

## A small excursion - Structured Programming

Control flow should fall into one of four patterns:

- Sequence: One block is executed after another
- Selection: One or more are executed (i.e. an `if`-statement)
- Iteration: A block is executed more than once (i.e. a loop)
- Recursion: A block calls itself until an exit condition is met (i.e. recursion!)

Do you know any keyword in C which *doesn't* quite adhere to that but can instead totally spaghetti your control flow? `goto`

# Thread models - Many To One

Do you know a programming language / runtime using that?

E.g. nodejs using its „event loop“

## A small excursion - Structured Programming

Control flow should fall into one of four patterns:

- Sequence: One block is executed after another
- Selection: One or more are executed (i.e. an `if`-statement)
- Iteration: A block is executed more than once (i.e. a loop)
- Recursion: A block calls itself until an exit condition is met (i.e. recursion!)

Do you know any keyword in C which *doesn't* quite adhere to that but can instead totally spaghetti your control flow? `goto`

Famous paper by a proponent of Structured Programming:

„*Go To Statement Considered Harmful*“ by Edsger W. Dijkstra

### And what about threads?

- Can outlive the methods they were spawned in

### And what about threads?

- Can outlive the methods they were spawned in
- Can use variables and fields after they went out of scope in a method

### And what about threads?

- Can outlive the methods they were spawned in
- Can use variables and fields after they went out of scope in a method
- Can split up or transfer their control flow arbitrarily

### And what about threads?

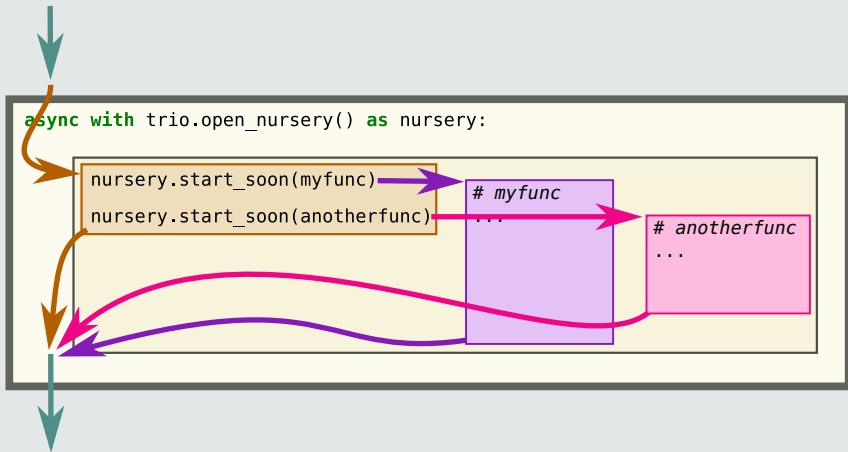
- Can outlive the methods they were spawned in
- Can use variables and fields after they went out of scope in a method
- Can split up or transfer their control flow arbitrarily

So that might sound familiar...



# Thread models - Many To One

## Structured Concurrency



Taken from [vorpus.org](https://vorpus.org)

Nice, but what does this have to do with ULTs?

### Nice, but what does this have to do with ULTs?

- Spawning lots of threads for small operations *is too slow otherwise*

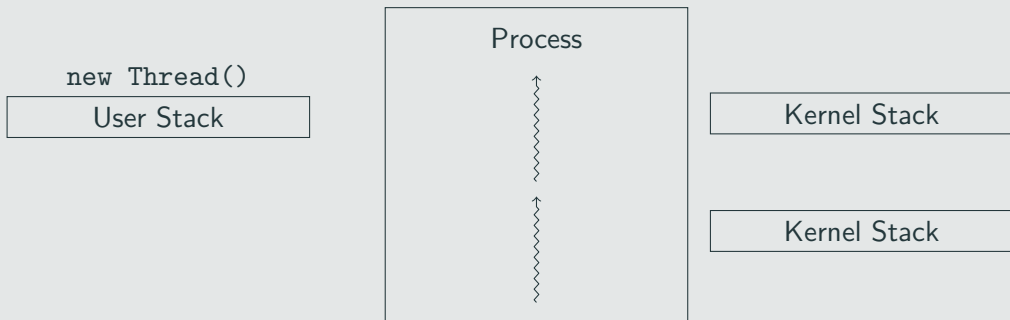
Further reading:

[Notes on Structured Concurrency](#)

[ULTs and Structured concurrency in Java - Project Loom](#)

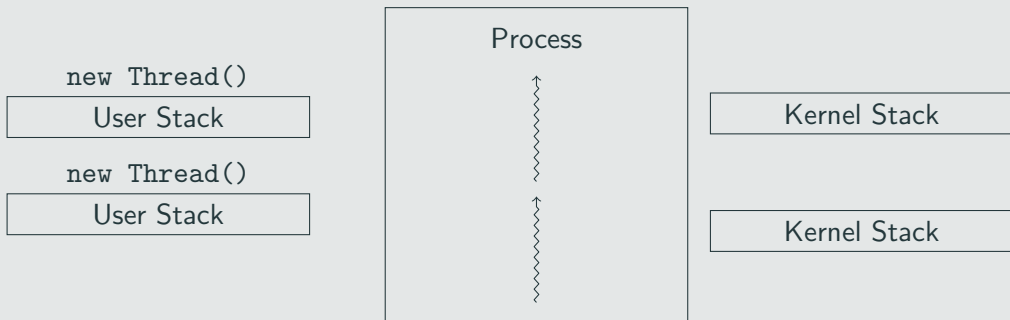
# Thread models - Many To Many

## Many To Many




# Thread models - Many To Many

## Many To Many




# Thread models - Many To Many

## Many To Many




# Thread models - Many To Many

## Many To Many



# Thread models - Many To Many


## Many To Many





# Thread models - Many To Many

## Many To Many



### Problems and benefits of *Many To Many*?

Important: The kernel *knows about the user level scheduler*

### Problems and benefits of *Many To Many*?

Important: The kernel *knows about the user level scheduler*

+ **Scales with core count?**

### Problems and benefits of *Many To Many*?

Important: The kernel *knows about the user level scheduler*

- + Scales with core count

### Problems and benefits of *Many To Many*?

Important: The kernel *knows about the user level scheduler*

- + Scales with core count
- + **Conceptually easy** — the OS does the hard stuff?

### Problems and benefits of *Many To Many*?

Important: The kernel *knows about the user level scheduler*

- + Scales with core count
- Harder to implement — the OS doesn't help you much

### Problems and benefits of *Many To Many*?

Important: The kernel *knows about the user level scheduler*

- + Scales with core count
- Harder to implement — the OS doesn't help you much
- + **Blocking does not affect other threads?**

## Problems and benefits of *Many To Many*?

Important: The kernel *knows about the user level scheduler*

- + Scales with core count
- Harder to implement — the OS doesn't help you much
- + Blocking does not affect other threads as the kernel informs the user scheduler (upcalls) and does *not* pause the whole kernel level thread



## Problems and benefits of *Many To Many*?

Important: The kernel *knows about the user level scheduler*

- + Scales with core count
- Harder to implement — the OS doesn't help you much
- + Blocking does not affect other threads as the kernel informs the user scheduler (upcalls) and does *not* pause the whole kernel level thread
- + **Can piggy-back onto the OS scheduler?**

## Problems and benefits of *Many To Many*?

Important: The kernel *knows about the user level scheduler*

- + Scales with core count
- Harder to implement — the OS doesn't help you much
- + Blocking does not affect other threads as the kernel informs the user scheduler (upcalls) and does *not* pause the whole kernel level thread
- Can *not* piggy-back onto the OS scheduler

## Problems and benefits of *Many To Many*?

Important: The kernel *knows about the user level scheduler*

- + Scales with core count
- Harder to implement — the OS doesn't help you much
- + Blocking does not affect other threads as the kernel informs the user scheduler (upcalls) and does *not* pause the whole kernel level thread
- Can *not* piggy-back onto the OS scheduler
- **Must** piggy-back onto the OS scheduler?

## Problems and benefits of *Many To Many*?

Important: The kernel *knows about the user level scheduler*

- + Scales with core count
- Harder to implement — the OS doesn't help you much
- + Blocking does not affect other threads as the kernel informs the user scheduler (upcalls) and does *not* pause the whole kernel level thread
- Can *not* piggy-back onto the OS scheduler
- + Can implement its **own** scheduler

## Problems and benefits of *Many To Many*?

Important: The kernel *knows about the user level scheduler*

- + Scales with core count
- Harder to implement — the OS doesn't help you much
- + Blocking does not affect other threads as the kernel informs the user scheduler (upcalls) and does *not* pause the whole kernel level thread
- Can *not* piggy-back onto the OS scheduler
- + Can implement its **own** scheduler
- **Relatively high overhead due to context switches?**

## Problems and benefits of *Many To Many*?

Important: The kernel *knows about the user level scheduler*

- + Scales with core count
- Harder to implement — the OS doesn't help you much
- + Blocking does not affect other threads as the kernel informs the user scheduler (upcalls) and does *not* pause the whole kernel level thread
- Can *not* piggy-back onto the OS scheduler
- + Can implement its **own** scheduler
- + Low overhead during context switches (unless you need to interact with the kernel threads, e.g. to schedule between them)

## Problems and benefits of *Many To Many*?

Important: The kernel *knows about the user level scheduler*

- + Scales with core count
- Harder to implement — the OS doesn't help you much
- + Blocking does not affect other threads as the kernel informs the user scheduler (upcalls) and does *not* pause the whole kernel level thread
- Can *not* piggy-back onto the OS scheduler
- + Can implement its **own** scheduler
- + Low overhead during context switches (unless you need to interact with the kernel threads, e.g. to schedule between them)
- **Relatively high overhead *when creating one?***

## Problems and benefits of *Many To Many*?

Important: The kernel *knows about the user level scheduler*


- + Scales with core count
- Harder to implement — the OS doesn't help you much
- + Blocking does not affect other threads as the kernel informs the user scheduler (upcalls) and does *not* pause the whole kernel level thread
- Can *not* piggy-back onto the OS scheduler
- + Can implement its **own** scheduler
- + Low overhead during context switches (unless you need to interact with the kernel threads, e.g. to schedule between them)
- + Low overhead when creating one



# Thread models - Many To Many

## Problems the second

You have all attended SWT 1! So let's have a look.




And preemption is now possible, which might complicate user code.

# Thread models - Many To Many

## Problems the second

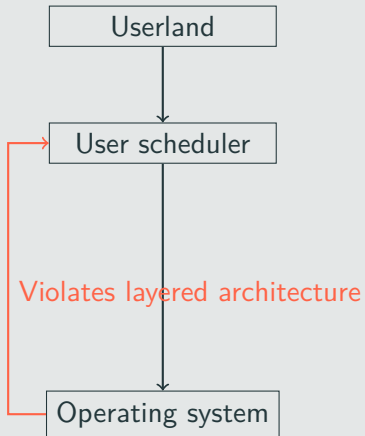
You have all attended SWT 1! So let's have a look.



# Thread models - Many To Many

## Problems the second


You have all attended SWT 1! So let's have a look.



# Thread models - Many To Many

## Problems the second

You have all attended SWT 1! So let's have a look.



And preemption is now possible, which might complicate user code.

**What events can trigger a context switch?**

### What events can trigger a context switch?

- Voluntary:

### What events can trigger a context switch?

- Voluntary: yield, blocking system call

### What events can trigger a context switch?

- Voluntary: yield, blocking system call
- Involuntary:
  - interrupts, exceptions, syscalls



## What events can trigger a context switch?

- Voluntary: yield, blocking system call
- Involuntary:
  - interrupts, exceptions, syscalls
  - end of time-slice

### What events can trigger a context switch?

- Voluntary: yield, blocking system call
- Involuntary:
  - interrupts, exceptions, syscalls
  - end of time-slice
  - high priority thread becoming ready

What events can trigger a context switch?

### What events can trigger a context switch?

- Most libraries only support *cooperative scheduling*

### What events can trigger a context switch?

- Most libraries only support *cooperative scheduling*
- Why is switching with preemption, interrupts, blocking system calls hard?

### What events can trigger a context switch?

- Most libraries only support *cooperative scheduling*
- Why is switching with preemption, interrupts, blocking system calls hard?  
Kernel is not aware of the ULTs and will return where it came from — *but not call out to the scheduler and carry on*
- The benefit of platforms: How can Java (using Project Loom) or Node.js switch on most of the above?

### What events can trigger a context switch?

- Most libraries only support *cooperative scheduling*
- Why is switching with preemption, interrupts, blocking system calls hard?  
Kernel is not aware of the ULTs and will return where it came from — *but not call out to the scheduler and carry on*
- The benefit of platforms: How can Java (using Project Loom) or Node.js switch on most of the above?  
You can not execute syscalls directly, but need to call library methods! Suspension points can be inserted there.

„Jobs are either I/O-bound or compute-bound. In neither case would user-level threads be a win. Why would one go for pure user-level threads at all?“



„Jobs are either I/O-bound or compute-bound. In neither case would user-level threads be a win. Why would one go for pure user-level threads at all?“

- Program structure (e.g. Structured Concurrency, channels or just easier pipelines)


**„Jobs are either I/O-bound or compute-bound. In neither case would user-level threads be a win. Why would one go for pure user-level threads at all?“**

- Program structure (e.g. Structured Concurrency, channels or just easier pipelines)
- The same or higher I/O throughput if on an abstracted platform

Should a fork do this?



Should a fork do this?



### Who is aware of the fork?

- The thread that executed the fork and the child

### Who is aware of the fork?

- The thread that executed the fork and the child
- Who else?

### Who is aware of the fork?

- The thread that executed the fork and the child
- Who else? Nobody!

### Who is aware of the fork?

- The thread that executed the fork and the child
- Who else? Nobody!

### Can you foresee any problems?

- If you copy the thread it might do weird things



### Who is aware of the fork?

- The thread that executed the fork and the child
- Who else? Nobody!

### Can you foresee any problems?

- If you copy the thread it might do weird things
- Locks are memory and will be copied as well! What if a thread holds a critical lock when it is forked?

### Who is aware of the fork?

- The thread that executed the fork and the child
- Who else? Nobody!

### Can you foresee any problems?

- If you copy the thread it might do weird things
- Locks are memory and will be copied as well! What if a thread holds a critical lock when it is forked?

### Do you need thread duplication for our shell example?

### Who is aware of the fork?

- The thread that executed the fork and the child
- Who else? Nobody!

### Can you foresee any problems?

- If you copy the thread it might do weird things
- Locks are memory and will be copied as well! What if a thread holds a critical lock when it is forked?

### Do you need thread duplication for our shell example?

No, we exec anyways.

## Who is aware of the fork?

- The thread that executed the fork and the child
- Who else? Nobody!

## Can you foresee any problems?

- If you copy the thread it might do weird things
- Locks are memory and will be copied as well! What if a thread holds a critical lock when it is forked?

## Do you need thread duplication for our shell example?

No, we exec anyways.

## Summary

`fork` is not as simple as it once was. Is it still a good abstraction?

What is the difference?

### What is the difference?

- Kernel Level Thread: The kernel knows about and manages the thread

### What is the difference?

- Kernel Level Thread: The kernel knows about and manages the thread
- Kernel Mode Thread: The thread runs in *kernel mode*

### What is the difference?

- Kernel Level Thread: The kernel knows about and manages the thread
- Kernel Mode Thread: The thread runs in *kernel mode*

### Why would you need kernel *mode* threads?



### What is the difference?

- Kernel Level Thread: The kernel knows about and manages the thread
- Kernel Mode Thread: The thread runs in *kernel mode*

### Why would you need kernel *mode* threads?

- Sometimes the OS needs to do some housekeeping that doesn't belong in a system call

### What is the difference?

- Kernel Level Thread: The kernel knows about and manages the thread
- Kernel Mode Thread: The thread runs in *kernel mode*

### Why would you need kernel *mode* threads?

- Sometimes the OS needs to do some housekeeping that doesn't belong in a system call
- Free some pages, swap something in and out of memory

### What is the difference?

- Kernel Level Thread: The kernel knows about and manages the thread
- Kernel Mode Thread: The thread runs in *kernel mode*

### Why would you need kernel *mode* threads?

- Sometimes the OS needs to do some housekeeping that doesn't belong in a system call
- Free some pages, swap something in and out of memory
- Flush pages from the disk cache to the hard disk

### What is the difference?

- Kernel Level Thread: The kernel knows about and manages the thread
- Kernel Mode Thread: The thread runs in *kernel mode*

### Why would you need kernel *mode* threads?

- Sometimes the OS needs to do some housekeeping that doesn't belong in a system call
- Free some pages, swap something in and out of memory
- Flush pages from the disk cache to the hard disk
- Perform VFS Checkpointing

### What is the difference?

- Kernel Level Thread: The kernel knows about and manages the thread
- Kernel Mode Thread: The thread runs in *kernel mode*

### Why would you need kernel *mode* threads?

- Sometimes the OS needs to do some housekeeping that doesn't belong in a system call
- Free some pages, swap something in and out of memory
- Flush pages from the disk cache to the hard disk
- Perform VFS Checkpointing
- ...