

Betriebssysteme

6. Tutorium - Paging

Peter Bohner

6. Dezember 2023

ITEC - Operating Systems Group

- Danke an die 2, die regelmäßig abgeben :)
- Niemand hat die Fadenschwimmbecken-Aufgabe¹ gelöst

¹Tichy lässt grüßen

- Danke an die 2, die regelmäßig abgeben :)
- Niemand hat die Fadenschwimmbecken-Aufgabe¹ gelöst
- Allgemein: Nur genau die Frage beantworten, die in der Aufgabe gestellt wurde → Punktabzug bei Fehler

¹Tichy lässt grüßen

- Danke an die 2, die regelmäßig abgeben :)
- Niemand hat die Fadenschwimmbecken-Aufgabe¹ gelöst
- Allgemein: Nur genau die Frage beantworten, die in der Aufgabe gestellt wurde → Punktabzug bei Fehler
- Segmentierung: Register

¹Tichy lässt grüßen

- Danke an die 2, die regelmäßig abgeben :)
- Niemand hat die Fadenschwimmbecken-Aufgabe¹ gelöst
- Allgemein: Nur genau die Frage beantworten, die in der Aufgabe gestellt wurde → Punktabzug bei Fehler
- Segmentierung: Register STBR, STLR

¹Tichy lässt grüßen

- Danke an die 2, die regelmäßig abgeben :)
- Niemand hat die Fadenschwimmbecken-Aufgabe¹ gelöst
- Allgemein: Nur genau die Frage beantworten, die in der Aufgabe gestellt wurde → Punktabzug bei Fehler
- Segmentierung: Register STBR, STLR
- *memcpy* vs *memmove*

¹Tichy lässt grüßen

- Danke an die 2, die regelmäßig abgeben :)
- Niemand hat die Fadenschwimmbecken-Aufgabe¹ gelöst
- Allgemein: Nur genau die Frage beantworten, die in der Aufgabe gestellt wurde → Punktabzug bei Fehler
- Segmentierung: Register STBR, STLR
- *memcpy* vs *memmove*
- PThreads: Wie erstellt man einen Thread? → Wiederholen?

¹Tichy lässt grüßen

- Danke an die 2, die regelmäßig abgeben :)
- Niemand hat die Fadenschwimmbecken-Aufgabe¹ gelöst
- Allgemein: Nur genau die Frage beantworten, die in der Aufgabe gestellt wurde → Punktabzug bei Fehler
- Segmentierung: Register STBR, STLR
- *memcpy* vs *memmove*
- PThreads: Wie erstellt man einen Thread? → Wiederholen?
- Fragen zum ÜB?

¹Tichy lässt grüßen

- Danke an die 2, die regelmäßig abgeben :)
- Niemand hat die Fadenschwimmbecken-Aufgabe¹ gelöst
- Allgemein: Nur genau die Frage beantworten, die in der Aufgabe gestellt wurde → Punktabzug bei Fehler
- Segmentierung: Register STBR, STLR
- *memcpy* vs *memmove*
- PThreads: Wie erstellt man einen Thread? → Wiederholen?
- Fragen zum ÜB?
- Macht jemand Advent of Code?

¹Tichy lässt grüßen

- Danke an die 2, die regelmäßig abgeben :)
- Niemand hat die Fadenschwimmbecken-Aufgabe¹ gelöst
- Allgemein: Nur genau die Frage beantworten, die in der Aufgabe gestellt wurde → Punktabzug bei Fehler
- Segmentierung: Register STBR, STLR
- *memcpy* vs *memmove*
- PThreads: Wie erstellt man einen Thread? → Wiederholen?
- Fragen zum ÜB?
- Macht jemand Advent of Code?

¹Tichy lässt grüßen

Paging

What is that? What is the difference to Segmentation?

What is that? What is the difference to Segmentation?

- Virtual memory is broken into fixed-size chunks (**pages**)

What is that? What is the difference to Segmentation?

- Virtual memory is broken into fixed-size chunks (**pages**)
- Physical memory is broken into fixed-size chunks of the same size (**frames**)

What is that? What is the difference to Segmentation?

- Virtual memory is broken into fixed-size chunks (**pages**)
- Physical memory is broken into fixed-size chunks of the same size (**frames**)
- Virtual pages are mapped to page frames

What is that? What is the difference to Segmentation?

- Virtual memory is broken into fixed-size chunks (**pages**)
- Physical memory is broken into fixed-size chunks of the same size (**frames**)
- Virtual pages are mapped to page frames Always?

What is that? What is the difference to Segmentation?

- Virtual memory is broken into fixed-size chunks (**pages**)
- Physical memory is broken into fixed-size chunks of the same size (**frames**)
- Virtual pages are mapped to page frames Always?
No! (Paged out, zero pages, ...)

Benefits over Segmentation?

What is that? What is the difference to Segmentation?

- Virtual memory is broken into fixed-size chunks (**pages**)
- Physical memory is broken into fixed-size chunks of the same size (**frames**)
- Virtual pages are mapped to page frames Always?
No! (Paged out, zero pages, ...)

Benefits over Segmentation?

- Virtual memory does not need to map to *continuous* physical memory

What is that? What is the difference to Segmentation?

- Virtual memory is broken into fixed-size chunks (**pages**)
- Physical memory is broken into fixed-size chunks of the same size (**frames**)
- Virtual pages are mapped to page frames Always?
No! (Paged out, zero pages, ...)

Benefits over Segmentation?

- Virtual memory does not need to map to *continuous* physical memory
- Swapping in/out is easier

What is that? What is the difference to Segmentation?

- Virtual memory is broken into fixed-size chunks (**pages**)
- Physical memory is broken into fixed-size chunks of the same size (**frames**)
- Virtual pages are mapped to page frames Always?
No! (Paged out, zero pages, ...)

Benefits over Segmentation?

- Virtual memory does not need to map to *continuous* physical memory
- Swapping in/out is easier
- No external fragmentation, little internal

Segment and Page tables

Segment Number	Base	Limit
0	0xdead	0x00ef
1	0xf154	0x013a
2	0x0000	0x0000
3	0x0000	0x3fff

Paging - Single Level Page Table

Segment and Page tables

Virtual page number	Base	Limit
0	0xdead	0x00ef
1	0xf154	0x013a
2	0x0000	0x0000
3	0x0000	0x3fff

Segment and Page tables

Virtual page number	Frame number	Limit
0	0xdead	0x00ef
1	0xf154	0x013a
2	0x0000	0x0000
3	0x0000	0x3fff

Segment and Page tables

Virtual page number	Frame number	Limit
0	0xdead	0x00ef
1	0xf154	0x013a
2	0x0000	0x0000
3	0x0000	0x3fff

Paging - Single Level Page Table

Segment and Page tables

0x020123

Paging - Single Level Page Table

Segment and Page tables

0x020123

Paging - Single Level Page Table

Segment and Page tables

0x020123

Page Number

Offset

0x02	0x0123
------	--------

Paging - Single Level Page Table

Segment and Page tables

0x020123

Page Number

Offset

0x02


0x0123

pfn	flags
0xDE	
0xAD	
...	

Paging - Single Level Page Table

Segment and Page tables

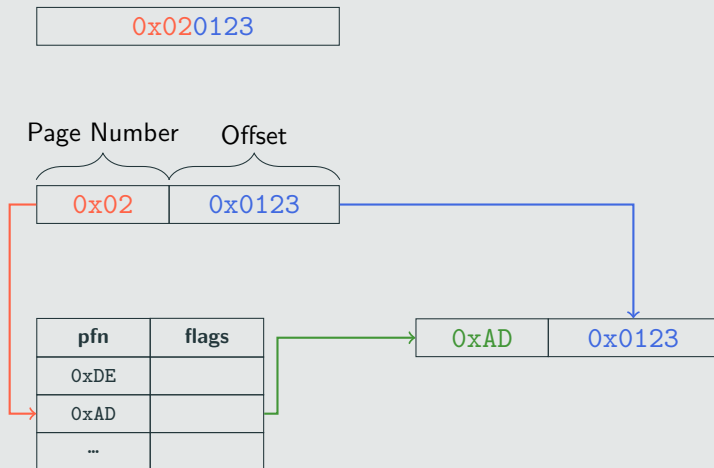
0x020123



pfn	flags
0xDE	
0xAD	
...	

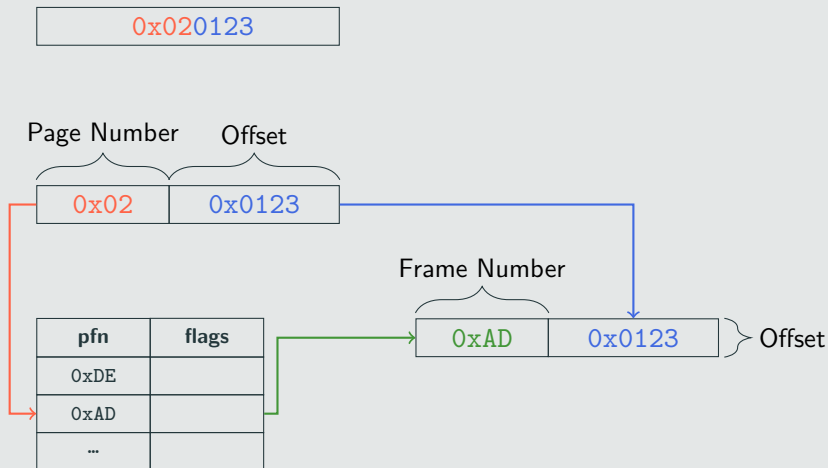
Paging - Single Level Page Table

Segment and Page tables



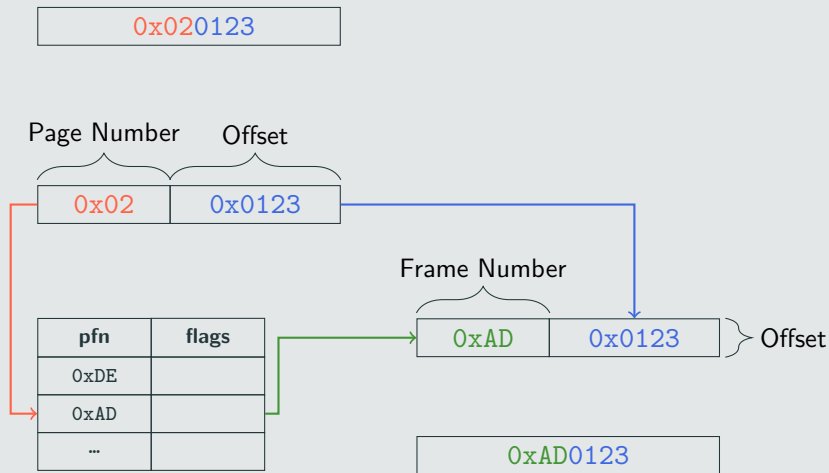
Paging - Single Level Page Table

Segment and Page tables



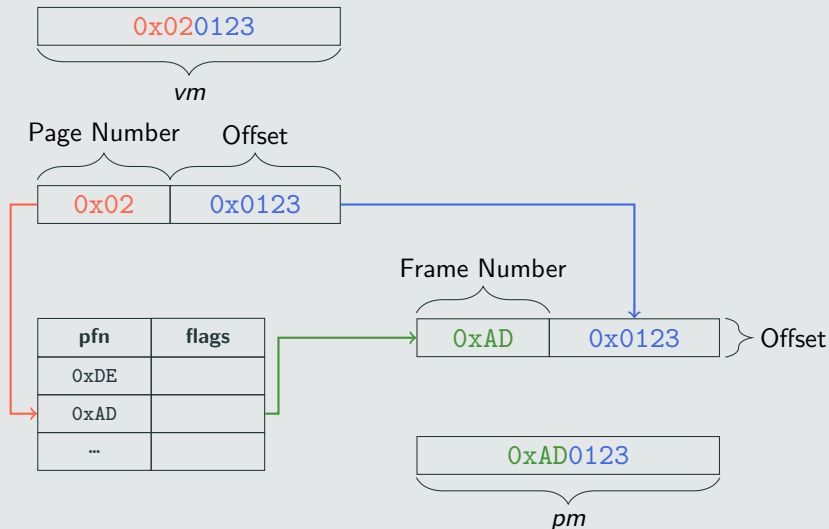
Paging - Single Level Page Table

Segment and Page tables



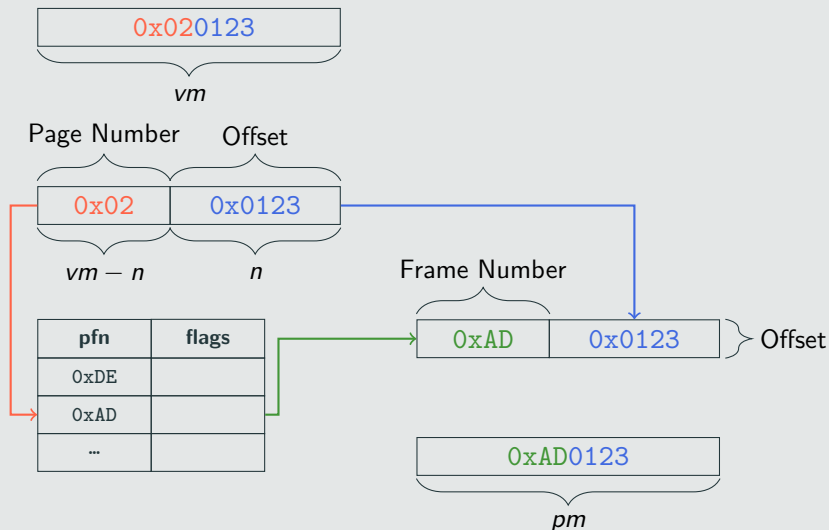
Paging - Single Level Page Table

Segment and Page tables



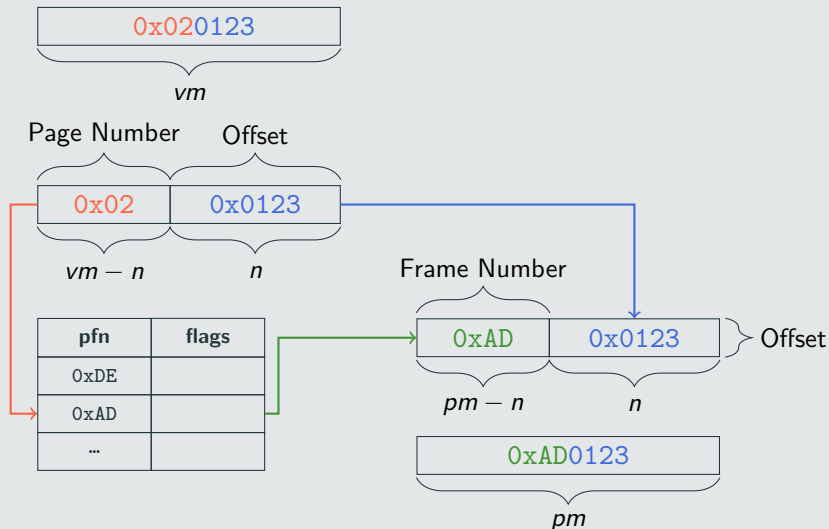
Paging - Single Level Page Table

Segment and Page tables



Paging - Single Level Page Table

Segment and Page tables



Single Level Page Table - Disadvantages

Why could that be a bit problematic?

Why could that be a bit problematic?

- Increases with the size of the size of the *virtual* address space

Why could that be a bit problematic?

- Increases with the size of the size of the *virtual* address space
- 64 Bit AS, 4KiB (2^{12}) pages $\Rightarrow n = 12 \Rightarrow$

Why could that be a bit problematic?

- Increases with the size of the size of the *virtual* address space
- 64 Bit AS, 4KiB (2^{12}) pages $\Rightarrow n = 12 \Rightarrow 2^{vm-n} = 2^{64-12} = 2^{52}$

Why could that be a bit problematic?

- Increases with the size of the size of the *virtual* address space
- 64 Bit AS, 4KiB (2^{12}) pages $\Rightarrow n = 12 \Rightarrow 2^{vm-n} = 2^{64-12} = 2^{52}$

\Rightarrow If every entry was 1 Bit we'd need (asking units...)

You have: 2^{52} bit

You want: tebibyte

* 512

/ 0.001953125

- You might *not* have that much memory to spare :)

Single Level Page Table - Disadvantages

Math is fun, let's do some math

Calculate the space requirements for a single level page table with

- 32-bit virtual addresses, 4KiB pages, 4 bytes per page table entry
- 48-bit virtual addresses, 4KiB pages, 4 bytes per page table entry

Single Level Page Table - Disadvantages

Math is fun, let's do some math

Calculate the space requirements for a single level page table with

- 32-bit virtual addresses, 4KiB pages, 4 bytes per page table entry
- 48-bit virtual addresses, 4KiB pages, 4 bytes per page table entry

32-bit

- $vm = 32, 4Kib = 2^{12} \Rightarrow n = 12$

Single Level Page Table - Disadvantages

Math is fun, let's do some math

Calculate the space requirements for a single level page table with

- 32-bit virtual addresses, 4KiB pages, 4 bytes per page table entry
- 48-bit virtual addresses, 4KiB pages, 4 bytes per page table entry

32-bit

- $vm = 32, 4Kib = 2^{12} \Rightarrow n = 12$
- $2^{32-12} = 2^{20}$ entries $\Rightarrow 2^{20} \cdot 2^2 = 2^{22}$ Byte (4 MiB)

Single Level Page Table - Disadvantages

Math is fun, let's do some math

Calculate the space requirements for a single level page table with

- 32-bit virtual addresses, 4KiB pages, 4 bytes per page table entry
- 48-bit virtual addresses, 4KiB pages, 4 bytes per page table entry

32-bit

- $vm = 32, 4Kib = 2^{12} \Rightarrow n = 12$
- $2^{32-12} = 2^{20}$ entries $\Rightarrow 2^{20} \cdot 2^2 = 2^{22}$ Byte (4 MiB)

48-bit

- $vm = 48, 4Kib = 2^{12} \Rightarrow n = 12$
- $2^{48-12} = 2^{36}$ entries $\Rightarrow 2^{36} \cdot 2^2 = 2^{38}$ Byte (256 GiB)

Alternatives to Single Level Page Tables



Mutli-Level page tables

Alternatives to Single Level Page Tables

Mutli-Level page tables


1st Level PT



0xDE	AD	BE	EF
------	----	----	----


Alternatives to Single Level Page Tables

Mutli-Level page tables




Alternatives to Single Level Page Tables

Mutli-Level page tables



Alternatives to Single Level Page Tables

Mutli-Level page tables




Benefits and Drawbacks?

- Pointer chasing down each level \Rightarrow More memory accesses

Alternatives to Single Level Page Tables

Mutli-Level page tables



Benefits and Drawbacks?

- Pointer chasing down each level \Rightarrow More memory accesses
- + Address spaces are *sparse* \Rightarrow Only instantiate page tables you need

What do those do?

What do those do?

- Map *physical* addresses to virtual ones

What do those do?

- Map *physical* addresses to virtual ones
- Why?

What do those do?

- Map *physical* addresses to virtual ones
- Why? Physical address space is much smaller and

What do those do?

- Map *physical* addresses to virtual ones
- Why? Physical address space is much smaller and you don't need a table per address space

What do those do?

- Map *physical* addresses to virtual ones
- Why? Physical address space is much smaller and you don't need a table per address space
- If you don't have a table per address space, how can you still have multiple?

What do those do?

- Map *physical* addresses to virtual ones
- Why? Physical address space is much smaller and you don't need a table per address space
- If you don't have a table per address space, how can you still have multiple?

⇒ e.g. Linked List of entries per physical frame

What do those do?

- Map *physical* addresses to virtual ones
- Why? Physical address space is much smaller and you don't need a table per address space
- If you don't have a table per address space, how can you still have multiple?

⇒ e.g. Linked List of entries per physical frame

Any drawbacks?

What do those do?

- Map *physical* addresses to virtual ones
 - Why? Physical address space is much smaller and you don't need a table per address space
 - If you don't have a table per address space, how can you still have multiple?
- ⇒ e.g. Linked List of entries per physical frame

Any drawbacks?

- We mostly care about *the other direction*, i.e. virtual ⇒ physical

What do those do?

- Map *physical* addresses to virtual ones
 - Why? Physical address space is much smaller and you don't need a table per address space
 - If you don't have a table per address space, how can you still have multiple?
- ⇒ e.g. Linked List of entries per physical frame

Any drawbacks?

- We mostly care about *the other direction*, i.e. virtual ⇒ physical
- That requires iteration :(

Inverted Page Tables

How can we speed them up?

After having attended *Algorithmen I* we all know:

Inverted Page Tables

How can we speed them up?

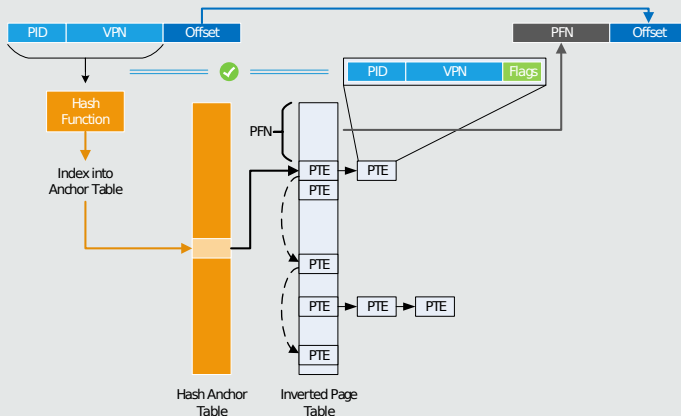
After having attended *Algorithmen I* we all know: Hacktables are $O(1)$

Inverted Page Tables

How can we speed them up?

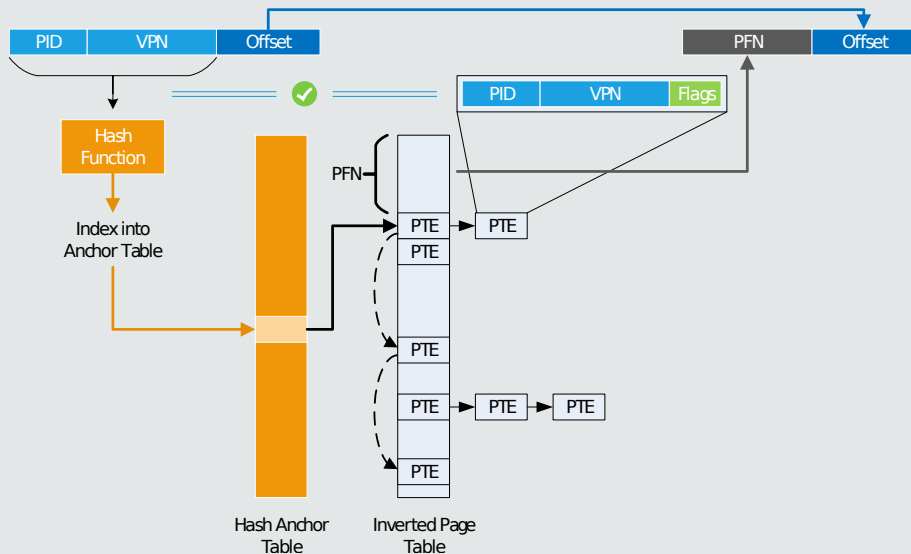
After having attended *Algorithmen I* we all know: Hacktables are $O(1)$

Hashed inverted page tables



Inverted Page Tables

Hashed inverted page tables



Why do pages have a valid/present bit?

What is it for?

Why do pages have a valid/present bit?

What is it for?

- The page might not be present. Why?

Why do pages have a valid/present bit?

What is it for?

- The page might not be present. Why? Swapped out,

Why do pages have a valid/present bit?

What is it for?

- The page might not be present. Why? Swapped out, not zeroed yet, ...

What happens on access if it is not set?

Why do pages have a valid/present bit?

What is it for?

- The page might not be present. Why? Swapped out, not zeroed yet, ...

What happens on access if it is not set?

- Page fault!
- Handle it and do sth. sensible (or crash the process...)

TLB



What does that refer to?

What does that refer to?

Translation **L**ookaside **B**uffer. What's that for?

What does that refer to?

Translation **L**ookaside **B**uffer. What's that for?

- Lookup from virtual to physical address can be *slow*

What does that refer to?

Translation Lookaside Buffer. What's that for?

- Lookup from virtual to physical address can be *slow*
- You need to translate addresses *all the time*

What does that refer to?

Translation Lookaside Buffer. What's that for?

- Lookup from virtual to physical address can be *slow*
- You need to translate addresses *all the time*

⇒ There is no problem you can't solve with another caching layer
(except having too many caching layers) [Nearly the Fundamental theorem of software engineering](#)

TLB layout

p	offset
---	--------


TLB layout

p	offset
---	--------

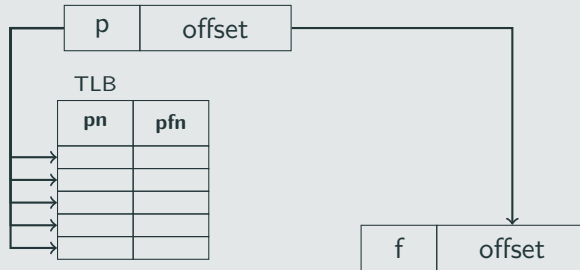
TLB

pn	pfn

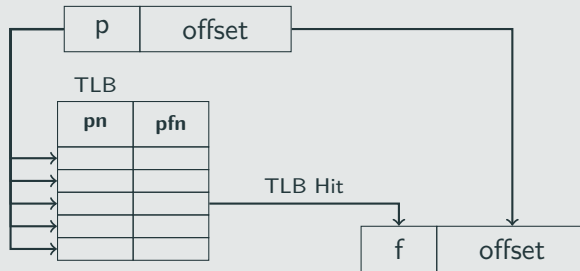
TLB layout




TLB layout




TLB layout



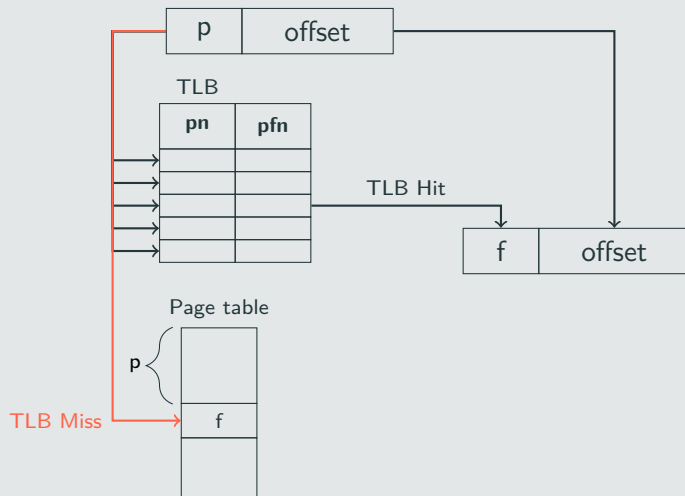
TLB layout



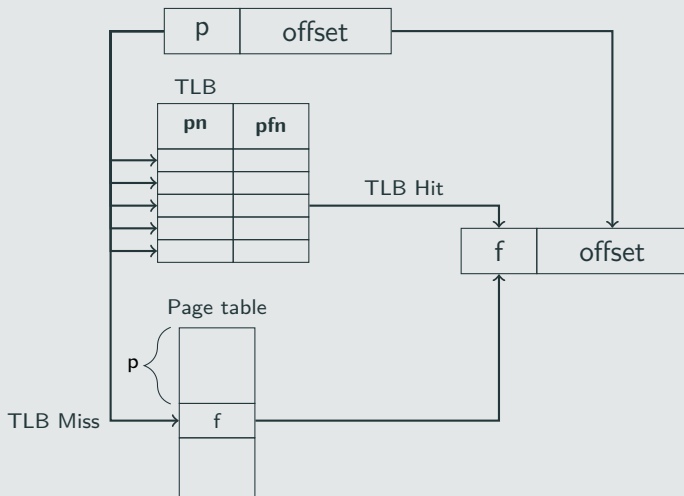
TLB layout



TLB layout



TLB layout



What is the difference between software and hardware walked TLBs?

What is the difference between software and hardware walked TLBs?

The TLB is *the same*, what differs is the behaviour on a *cache miss*.

What is the difference between software and hardware walked TLBs?

The TLB is *the same*, what differs is the behaviour on a *cache miss*.

Hardware walked

- Hardware looks at your page table
- Hardware *resolves* the physical address using it

What is the difference between software and hardware walked TLBs?

The TLB is *the same*, what differs is the behaviour on a *cache miss*.

Hardware walked

- Hardware looks at your page table
 - Hardware *resolves* the physical address using it
 - On failure a page fault is raised
- ⇒ What does that imply?

What is the difference between software and hardware walked TLBs?

The TLB is *the same*, what differs is the behaviour on a *cache miss*.

Hardware walked

- Hardware looks at your page table
- Hardware *resolves* the physical address using it
- On failure a page fault is raised

⇒ What does that imply? Page table layout is **fixed**

What is the difference between software and hardware walked TLBs?

The TLB is *the same*, what differs is the behaviour on a *cache miss*.

Hardware walked

- Hardware looks at your page table
- Hardware *resolves* the physical address using it
- On failure a page fault is raised

⇒ What does that imply? Page table layout is **fixed**

Software walked

- Transfers control to TLB miss handler
- Kernel routine walks the page table and tries to find a mapping

What is the difference between software and hardware walked TLBs?

The TLB is *the same*, what differs is the behaviour on a *cache miss*.

Hardware walked

- Hardware looks at your page table
- Hardware *resolves* the physical address using it
- On failure a page fault is raised

⇒ What does that imply? Page table layout is **fixed**

Software walked

- Transfers control to TLB miss handler
- Kernel routine walks the page table and tries to find a mapping
- Loads that mapping into the TLB and can choose which entry to evict!
- If there is none ⇒ Jump to page fault handler

What are benefits / drawbacks of software walked TLBs?

What are benefits / drawbacks of software walked TLBs?

- + Free to choose page table layout (fit your algorithm)

What are benefits / drawbacks of software walked TLBs?

- + Free to choose page table layout (fit your algorithm)
- + Free to choose which TLB entries to evict \Rightarrow Use optimized strategy

What are benefits / drawbacks of software walked TLBs?

- + Free to choose page table layout (fit your algorithm)
- + Free to choose which TLB entries to evict \Rightarrow Use optimized strategy
- Greater overhead

What information does the TLB store?

What information does the TLB store?

- Physical Frame Number, Virtual Page Number

What information does the TLB store?

- Physical Frame Number, Virtual Page Number
- Valid / Present Bit

What information does the TLB store?

- Physical Frame Number, Virtual Page Number
- Valid / Present Bit
- Modified bit, permissions, ...

When is a TLB miss or page fault raised?

Due to the TLB:

- Software walked: No entry found in TLB \Rightarrow TLB miss exception

When is a TLB miss or page fault raised?

Due to the TLB:

- Software walked: No entry found in TLB \Rightarrow TLB miss exception
- Hardware walked: No mapping in page table \Rightarrow Page fault

And other problems:

When is a TLB miss or page fault raised?

Due to the TLB:

- Software walked: No entry found in TLB \Rightarrow TLB miss exception
- Hardware walked: No mapping in page table \Rightarrow Page fault

And other problems: *Invalid access to a mapped page*

When is a TLB miss or page fault raised?

Due to the TLB:

- Software walked: No entry found in TLB \Rightarrow TLB miss exception
- Hardware walked: No mapping in page table \Rightarrow Page fault

And other problems: *Invalid access to a mapped page*

- Software walked: Raise some kind of TLB fault if the page is e.g. read-only

When is a TLB miss or page fault raised?

Due to the TLB:

- Software walked: No entry found in TLB \Rightarrow TLB miss exception
- Hardware walked: No mapping in page table \Rightarrow Page fault

And other problems: *Invalid access to a mapped page*

- Software walked: Raise some kind of TLB fault if the page is e.g. read-only
- Hardware walked: Page fault raised, page fault handler has to find out what happend

Page Fault Handling

What is Demand-Paging and Pre-Paging?

Demand-Paging:

What is Demand-Paging and Pre-Paging?

Demand-Paging:

- Load pages on-demand, right when they are needed

What is Demand-Paging and Pre-Paging?

Demand-Paging:

- Load pages on-demand, right when they are needed

Pre-Paging:

- Loaded Pages speculatively in batches, even *before* you need them

What is Demand-Paging and Pre-Paging?

Demand-Paging:

- Load pages on-demand, right when they are needed

Pre-Paging:

- Loaded Pages speculatively in batches, even *before* you need them

Why would you (not?) use Demand-Paging?

What is Demand-Paging and Pre-Paging?

Demand-Paging:

- Load pages on-demand, right when they are needed

Pre-Paging:

- Loaded Pages speculatively in batches, even *before* you need them

Why would you (not?) use Demand-Paging?

- + Only loads needed data \Rightarrow Less memory wasted

What is Demand-Paging and Pre-Paging?

Demand-Paging:

- Load pages on-demand, right when they are needed

Pre-Paging:

- Loaded Pages speculatively in batches, even *before* you need them

Why would you (not?) use Demand-Paging?

- + Only loads needed data \Rightarrow Less memory wasted
- Generates lots of page faults before working set is in memory

Why would you (not?) use Pre-Paging?

Why would you (not?) use Pre-Paging?

- + Might reduce number of page faults

Why would you (not?) use Pre-Paging?

- + Might reduce number of page faults
- Loads more than needed \Rightarrow Wasteful

Why would you (not?) use Pre-Paging?

- + Might reduce number of page faults
- Loads more than needed \Rightarrow Wasteful
- More I/O \Rightarrow Slower?

Why would you (not?) use Pre-Paging?

- + Might reduce number of page faults
- Loads more than needed \Rightarrow Wasteful
- More I/O \Rightarrow Slower?
- + HDDs a lot faster when reading chunks

Different kind of page faults

Not all pages are created equal. Do you have any idea what types of page faults typically exist?

On-Demand Paging

Different kind of page faults

Not all pages are created equal. Do you have any idea what types of page faults typically exist?



Kernel Space

The diagram consists of a solid purple rectangle labeled 'Kernel Space'. Below it is a dashed-line rectangle labeled 'No access'.

No access

On-Demand Paging

Different kind of page faults

Not all pages are created equal. Do you have any idea what types of page faults typically exist?

Kernel Space
No access
Buffer
No access
Stack

On-Demand Paging

Different kind of page faults


Not all pages are created equal. Do you have any idea what types of page faults typically exist?

Kernel Space
No access
Buffer
No access
Stack
No access
Heap

On-Demand Paging

Different kind of page faults

Not all pages are created equal. Do you have any idea what types of page faults typically exist?



On-Demand Paging

Different kind of page faults


Not all pages are created equal. Do you have any idea what types of page faults typically exist?

Kernel Space
No access
Buffer
No access
Stack
No access
Heap
No access
/root/some-file
No access

On-Demand Paging

Different kind of page faults

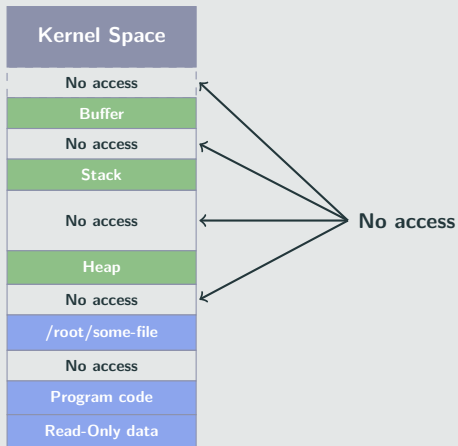
Not all pages are created equal. Do you have any idea what types of page faults typically exist?



On-Demand Paging

Different kind of page faults

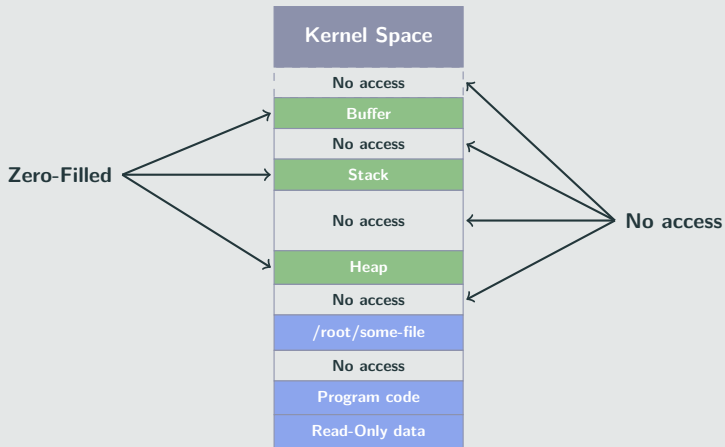
Not all pages are created equal. Do you have any idea what types of page faults typically exist?



On-Demand Paging

Different kind of page faults

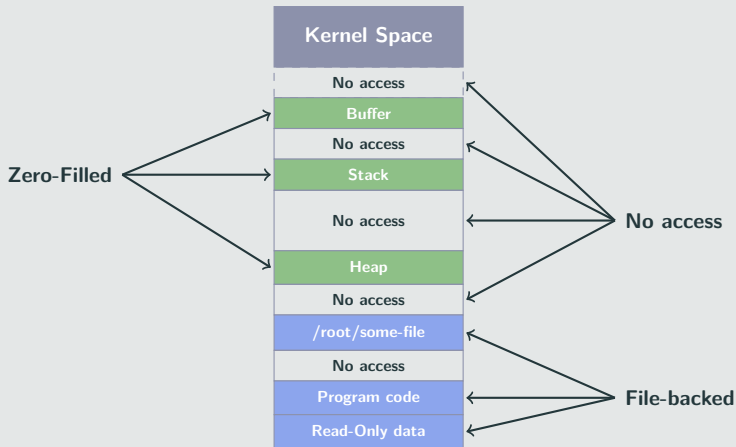
Not all pages are created equal. Do you have any idea what types of page faults typically exist?



On-Demand Paging

Different kind of page faults

Not all pages are created equal. Do you have any idea what types of page faults typically exist?



Different kind of page faults

Why are pages to generic memory zero filled?

Different kind of page faults

Why are pages to generic memory zero filled? It could leak other processes' memory otherwise (called a [Covert Channel](#)).

Different kind of page faults

Why are pages to generic memory zero filled? It could leak other processes' memory otherwise (called a [Covert Channel](#)).

And there is one other kind of page fault...

Different kind of page faults

Why are pages to generic memory zero filled? It could leak other processes' memory otherwise (called a [Covert Channel](#)).

And there is one other kind of page fault...

...The page was stolen by the OS and swapped out!

Different kind of page faults

Why are pages to generic memory zero filled? It could leak other processes' memory otherwise (called a [Covert Channel](#)).

And there is one other kind of page fault...

...The page was stolen by the OS and swapped out!

Also supported on some systems: *Purgable memory*. Stolen from [Apple](#) and also implemented in [SerenityOS](#) in [this video](#).

What kind of information does the page fault handler need?

What kind of information does the page fault handler need?

- Access flags: Can the user perform the operation on this page?

What kind of information does the page fault handler need?

- Access flags: Can the user perform the operation on this page?
- Where to find the most recent version (different for zero filled, file backed, etc.)

How could you implement Copy-on-Write memory?

How could you implement Copy-on-Write memory?

- Mark memory as read-only on fork
- Add an additional CoW flag: When a page fault is raised check it, copy the page and clear the CoW and ro flag



XKCD 912 - Manual Override

F R A G E N ?



<https://forms.gle/9CwJSKidKibubran9>

Bis nächste Woche :)