

Betriebssysteme

9. Tutorium - IPC

Peter Bohner

9. Januar 2024

ITEC - Operating Systems Group

- Frohes Neues!

- Frohes Neues!
- Neujahrsvorsatz: OS-Übungsblätter bearbeiten

- Frohes Neues!
- Neujahrsvorsatz: OS-Übungsblätter bearbeiten
- `<rant>` Wenn man vorher explizit sagt das man kommt, hat man das auch zu tun, oder sich abzumelden! `</rant>`

Inter-Process-Communication

How can different processes interact?

How can different processes interact?

- Shared memory (implicitly

How can different processes interact?

- Shared memory (implicitly as they share an address space and explicitly via a shared memory area)

How can different processes interact?

- Shared memory (implicitly as they share an address space and explicitly via a shared memory area)
- OS facilities (e.g. messages, pipes, Signals, sockets)

How can different processes interact?

- Shared memory (implicitly as they share an address space and explicitly via a shared memory area)
- OS facilities (e.g. messages, pipes, Signals, sockets)
- High level abstractions (files, database entries)

What mechanism do you need for IPC in an imperfect world?

What mechanism do you need for IPC in an imperfect world?

Timeouts! You don't want to wait for buggy programs or poor dead ones :(

Asynchronous send Operations require a buffer. Where do you put that?

- Actually, *why do they need a buffer?*

Asynchronous send Operations require a buffer. Where do you put that?

- Actually, *why do they need a buffer?* Sent but not yet received messages

Asynchronous send Operations require a buffer. Where do you put that?

- Actually, *why do they need a buffer?* Sent but not yet received messages
- They can be in the *receiver's* address space, the *sender's* AS or in the kernel

Asynchronous send Operations require a buffer. Where do you put that?

- Actually, *why do they need a buffer?* Sent but not yet received messages
- They can be in the *receiver's* address space, the *sender's* AS or in the kernel

Buffering in the Receiver's Address Space

- Is that a good idea?

Asynchronous send Operations require a buffer. Where do you put that?

- Actually, *why do they need a buffer?* Sent but not yet received messages
- They can be in the *receiver's* address space, the *sender's* AS or in the kernel

Buffering in the Receiver's Address Space

- Is that a good idea? What happens when you have many clients?

Asynchronous send Operations require a buffer. Where do you put that?

- Actually, *why do they need a buffer?* Sent but not yet received messages
- They can be in the *receiver's* address space, the *sender's* AS or in the kernel

Buffering in the Receiver's Address Space

- Is that a good idea? What happens when you have many clients?
- ⇒ Does not scale well. You either allow every sender to allocate memory for you (*ouch*) or you might run out with many clients

Buffering in the Kernel

- Is that a good idea?

Buffering in the Kernel

- Is that a good idea? What happens when you have many clients?

Buffering in the Kernel

- Is that a good idea? What happens when you have many clients?
- ⇒ Does not scale well. You either allow every sender to allocate kernel memory (*ouch*) or you might run out with many clients

Buffering in the Kernel

- Is that a good idea? What happens when you have many clients?
- ⇒ Does not scale well. You either allow every sender to allocate kernel memory (*ouch*) or you might run out with many clients

Buffering in the Sender's Address Space

- Sounds good?

Buffering in the Kernel

- Is that a good idea? What happens when you have many clients?
- ⇒ Does not scale well. You either allow every sender to allocate kernel memory (*ouch*) or you might run out with many clients

Buffering in the Sender's Address Space

- Sounds good? Yea, we are running out of options.

Buffering in the Kernel

- Is that a good idea? What happens when you have many clients?
- ⇒ Does not scale well. You either allow every sender to allocate kernel memory (*ouch*) or you might run out with many clients

Buffering in the Sender's Address Space

- Sounds good? Yea, we are running out of options. But it's mostly alright!

Buffering in the Kernel

- Is that a good idea? What happens when you have many clients?
- ⇒ Does not scale well. You either allow every sender to allocate kernel memory (*ouch*) or you might run out with many clients

Buffering in the Sender's Address Space

- Sounds good? Yea, we are running out of options. But it's mostly alright!
- + We only need to store it once: The sender has it in a buffer somewhere anyways

Buffering in the Kernel

- Is that a good idea? What happens when you have many clients?
- ⇒ Does not scale well. You either allow every sender to allocate kernel memory (*ouch*) or you might run out with many clients

Buffering in the Sender's Address Space

- Sounds good? Yea, we are running out of options. But it's mostly alright!
- + We only need to store it once: The sender has it in a buffer somewhere anyways
- + Scales better, as each sender keeps their messages
- ± We need to tell the client when it can reclaim the buffer

You are a very popular process that receives and handles many messages.

- For simplicities sake you chose *Synchronous IPC*. What problem might occur with many clients?

The Joys of Synchronous IPC

You are a very popular process that receives and handles many messages.

- For simplicities sake you chose *Synchronous IPC*. What problem might occur with many clients?
- You spent your whole life *waiting for timeouts to expire*

You are a very popular process that receives and handles many messages.

- For simplicities sake you chose *Synchronous IPC*. What problem might occur with many clients?
- You spent your whole life *waiting for timeouts to expire*
- How could you solve that with a new syscall?

You are a very popular process that receives and handles many messages.

- For simplicities sake you chose *Synchronous IPC*. What problem might occur with many clients?
- You spent your whole life *waiting for timeouts to expire*
- How could you solve that with a new syscall? How does `send-and-receive`, which sends and instantly receives help?

The Joys of Synchronous IPC

You are a very popular process that receives and handles many messages.

- For simplicities sake you chose *Synchronous IPC*. What problem might occur with many clients?
- You spent your whole life *waiting for timeouts to expire*
- How could you solve that with a new syscall? How does `send-and-receive`, which sends and instantly receives help?
- The server can assume you are using it and set a zero timeout. After all, if you are using that syscall you *will* be waiting

How can you emulate asynchronous IPC using synchronous IPC?

How can you emulate asynchronous IPC using synchronous IPC?

- Use a Proxy thread

How can you emulate asynchronous IPC using synchronous IPC?

- Use a Proxy thread
 1. Copy message to proxy thread
 2. Proxy threads sends synchronously and might block until recipient calls receive

How can you emulate asynchronous IPC using synchronous IPC?

- Use a Proxy thread
 1. Copy message to proxy thread
 2. Proxy threads sends synchronously and might block until recipient calls receive

+ Allows async I/O

-

How can you emulate asynchronous IPC using synchronous IPC?

- Use a Proxy thread
 1. Copy message to proxy thread
 2. Proxy threads sends synchronously and might block until recipient calls receive
- + Allows async I/O
 - How many messages can you send?

How can you emulate asynchronous IPC using synchronous IPC?

- Use a Proxy thread
 1. Copy message to proxy thread
 2. Proxy threads sends synchronously and might block until recipient calls receive
- + Allows async I/O
 - How many messages can you send? Yea, one per thread...

How can you emulate Synchronous IPC using asynchronous IPC?

How can you emulate Synchronous IPC using asynchronous IPC?

What about

```
1  do {  
2      res = async_send(message);  
3  } while(res != MESSAGE_SENT);
```

How can you emulate Synchronous IPC using asynchronous IPC?

What about

```
1  do {  
2      res = async_send(message);  
3  } while(res != MESSAGE_SENT);
```

Will wait until it was *sent*, not until it was *received*.

How can you emulate Synchronous IPC using asynchronous IPC?

What about

```
1  do {  
2      res = async_send(message);  
3  } while(res != MESSAGE_SENT);
```

Will wait until it was *sent*, not until it was *received*.

⇒ Implement a very simple protocol involving multiple messages

How can you emulate Synchronous IPC using asynchronous IPC?

What about

```
1  do {  
2      res = async_send(message);  
3  } while(res != MESSAGE_SENT);
```

Will wait until it was *sent*, not until it was *received*.

⇒ Implement a very simple protocol involving multiple messages

⇒ Send ACK message on receiver side, wait for ACK to be received.

And receive?

How can you emulate Synchronous IPC using asynchronous IPC?

What about

```
1  do {  
2      res = async_send(message);  
3  } while(res != MESSAGE_SENT);
```

Will wait until it was *sent*, not until it was *received*.

⇒ Implement a very simple protocol involving multiple messages

⇒ Send ACK message on receiver side, wait for ACK to be received.

And receive?

Just loop until `async_receive` receives a message (that is not an ACK)

Critical Sections

Synchronizing Properly - Core Properties

Does this work?

```
1 void foo() { // called in parallel
2     if(random() < 0.5) {
3         sleep(10);
4     }
5     a += 10;
6 }
```

Synchronizing Properly - Core Properties

Does this work?

```
1 void foo() { // called in parallel
2     if(random() < 0.5) {
3         sleep(10);
4     }
5     a += 10;
6 }
```

Nope. Multiple threads can call `a += 10` at the same time.

Property 1: Mutual Exclusion

Synchronizing Properly - Notes on Terminology

Some boring definitions

```
1 void foo() { // called in parallel
2
3     // This is the code before the critical section
4     // ==> Entry section
5
6     // Here common data is accessed (e.g. shared variable).
7     // Only one thread might be in here at a time.
8     // ==> Critical Section
9     a += 10;
10
11    // This is the code after the critical section
12    // ==> Exit section
13 }
```

Everything else is the **Remainder section**

Synchronizing Properly - Core Properties

Does this work?

```
1 void foo() { // called in parallel
2     // The late bird catches the worm
3     while(me != selectLastWaiting()) {
4         sleep(10);
5     }
6     a += 10;
7 }
```


Synchronizing Properly - Core Properties

Does this work?

```
1 void foo() { // called in parallel
2     // The late bird catches the worm
3     while(me != selectLastWaiting()) {
4         sleep(10);
5     }
6     a += 10;
7 }
```

How long could a thread wait?

Synchronizing Properly - Core Properties

Does this work?

```
1 void foo() { // called in parallel
2     // The late bird catches the worm
3     while(me != selectLastWaiting()) {
4         sleep(10);
5     }
6     a += 10;
7 }
```

How long could a thread wait?

Forever :(

Property 2: Bounded Waiting

Does this work?

```
1 void foo() { // called in parallel
2     while(waitingThreadCount > 1) {
3         sleep(10);
4     }
5     a += 10;
6 }
```

Does this work?

```
1 void foo() { // called in parallel
2     while(waitingThreadCount > 1) {
3         sleep(10);
4     }
5     a += 10;
6 }
```

Threads *outside* the critical section prevent threads from *entering* it

⇒ There's no progress!

Property 3: Progress

And the last one isn't really a property of a correct solution

Property 4: **Performance**

Synchronizing Properly - Core Properties

Mutual Exclusion

Only *one thread* can enter the critical section at once.

Synchronizing Properly - Core Properties

Mutual Exclusion

Only *one thread* can enter the critical section at once.

Progress

Threads in the *remainder* section do not prevent threads from *entering* the critical section.

Synchronizing Properly - Core Properties

Mutual Exclusion

Only *one thread* can enter the critical section at once.

Progress

Threads in the *remainder* section do not prevent threads from *entering* the critical section.

Bounded Waiting

There is an upper bound on *how many* different threads can enter the CS while a thread is waiting.

Important: This is not a *time* bound!

Synchronizing Properly - Core Properties

Mutual Exclusion

Only *one thread* can enter the critical section at once.

Progress

Threads in the *remainder* section do not prevent threads from *entering* the critical section.

Bounded Waiting

There is an upper bound on *how many* different threads can enter the CS while a thread is waiting.

Important: This is not a *time* bound!

Performance

The time overhead of the synchronization primitive is low (for low / medium / high contention).

Synchronizing Properly - Core Properties

Does this code fulfill all properties?

```
1  /* aligned to cache lines */
2  volatile int next = 0;
3  volatile int executing = 0;
4
5  lock_acquire() {
6      /* atomic version of "ticket = next; next += 1;" */
7      int ticket = fetch_and_add(next, 1);
8      /* busy wait until the counter matches */
9      while (ticket != executing ) {}
10 }
11
12 lock_release() {
13     executing++;
14 }
```

Does this code fulfill all properties?

- Mutual Exclusion:

Does this code fulfill all properties?

- Mutual Exclusion: Yes (though if you have 2^{32} threads waiting it doesn't)
- Bounded Waiting:

Does this code fulfill all properties?

- Mutual Exclusion: Yes (though if you have 2^{32} threads waiting it doesn't)
- Bounded Waiting: Yes. Eventually my number is the next one!
- Progress:

Does this code fulfill all properties?

- Mutual Exclusion: Yes (though if you have 2^{32} threads waiting it doesn't)
- Bounded Waiting: Yes. Eventually my number is the next one!
- Progress: Yes, threads in the remainder section do not hinder any thread from entering the critical section. Only threads in the *entry section* can do that temporarily.

Remember the bank example?

Remember the bank example?

```
1  current = get_balance();  
2  current += delta; //  $\text{delta} \in \{-50, 100\}$   
3  set_balance(current);
```

How could you fix that?

Remember the bank example?

```
1  current = get_balance();  
2  current += delta; //  $\text{delta} \in \{-50, 100\}$   
3  set_balance(current);
```

How could you fix that?

Synchronize it!

Fixing The Bank

Remember the bank example?

```
1  current = get_balance();  
2  current += delta; // delta ∈ {−50,100}  
3  set_balance(current);
```

How could you fix that?

Synchronize it!

```
1  lock(L);  
2  current = get_balance();  
3  current += delta; // delta ∈ {−50,100}  
4  set_balance(current);  
5  unlock(L);
```

Race Conditions

The Parallel Wizard Strikes Again

```
1  current = get_balance();  
2  current += delta; // delta ∈ {−50, 100}  
3  set_balance(current);
```

What happens if this code is executed in *parallel* for the two values of `delta`?

The Parallel Wizard Strikes Again

```
1  current = get_balance();
2  current += delta; // delta ∈ {−50,100}
3  set_balance(current);
```

What happens if this code is executed in *parallel* for the two values of *delta*?

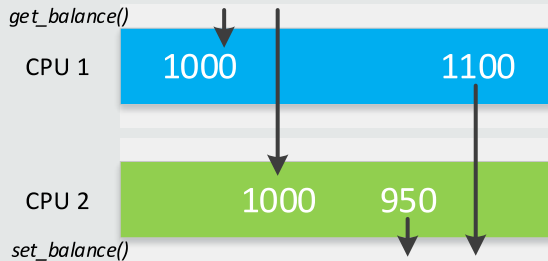
Unrolled

```
1  current = get_balance();
2  int tmp = current;
3  current = tmp + delta; // delta ∈ {−50,100}
4  set_balance(current);
```

Thread execution

Race Conditions

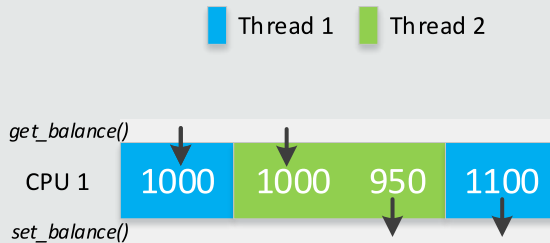
Thread execution



And if you run it concurrently on a single-core system?

Race Conditions

And if you run it concurrently on a single-core system?

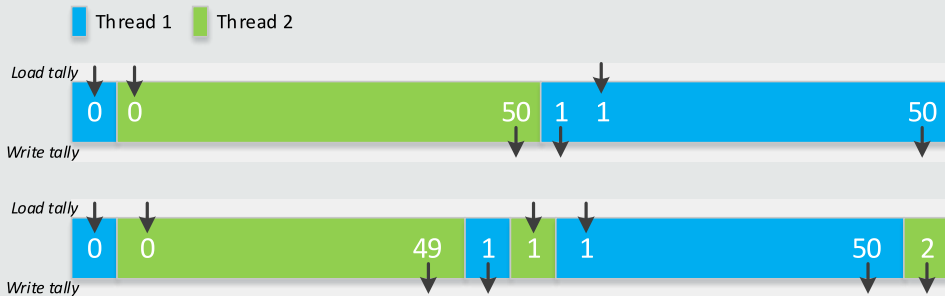


What is the value of tally at the end?

```
1 #include <stdio.h>
2
3 int tally;
4 void total(int N) {
5     for(int i = 0; i < N; i++)
6         tally += 1;
7 }
8
9 int main() {
10     tally = 0;
11
12     #pragma omp parallel for
13     for(int i = 0; i < 2; i++)
14         total(50);
15
16     printf("%d\n" , tally);
17     return 0;
18 }
```

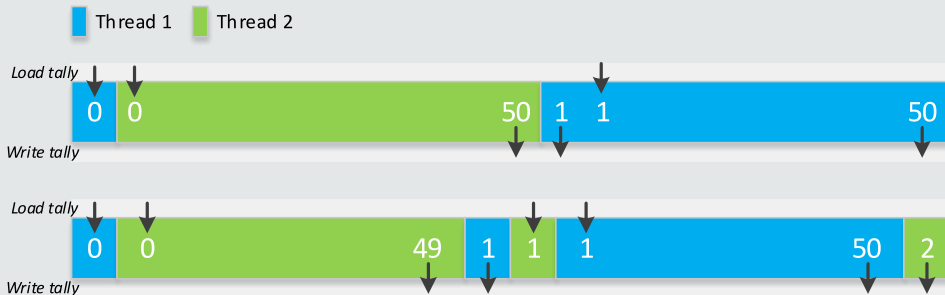
Having Fun With It...

This!



Having Fun With It...

This!

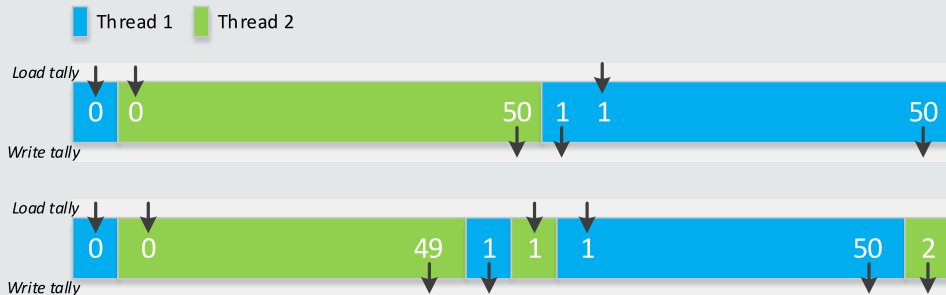


And what happens when we use $N, N > 0$ threads instead of 2?

The range is now

Having Fun With It...

This!



And what happens when we use $N, N > 0$ threads instead of 2?

The range is now $[2, 50 \cdot M]$. We can still have the same problem.

...And Adding ULTs!

What would happen if we used ULTs?

...And Adding ULTs!

What would happen if we used ULTs?

Can ULTs be pre-empted?

...And Adding ULTs!

What would happen if we used ULTs?

Can ULTs be pre-empted? With cooperatively-scheduled (the common case) no!

...And Adding ULTs!

What would happen if we used ULTs?

Can ULTs be pre-empted? With cooperatively-scheduled (the common case) no!

⇒ Everything will be fine, *unless* the threads yield *during the increment*

...And Adding ULTs!

What would happen if we used ULTs?

Can ULTs be pre-empted? With cooperatively-scheduled (the common case) no!

⇒ Everything will be fine, *unless* the threads yield *during the increment*

In that spirit: Does this work (with 1:1 threads)?

```
1 void total(int N) {  
2     for(int i = 0; i < N; ++ i) {  
3         tally += 1;  
4         sched_yield();  
5     }  
6 }
```

...And Adding ULTs!

What would happen if we used ULTs?

Can ULTs be pre-empted? With cooperatively-scheduled (the common case) no!

⇒ Everything will be fine, *unless* the threads yield *during the increment*

In that spirit: Does this work (with 1:1 threads)?

```
1 void total(int N) {  
2     for(int i = 0; i < N; ++ i) {  
3         tally += 1;  
4         sched_yield();  
5     }  
6 }
```

Nope. It might be less likely that it is interrupted during the increment (Why?

...And Adding ULTs!

What would happen if we used ULTs?

Can ULTs be pre-empted? With cooperatively-scheduled (the common case) no!

⇒ Everything will be fine, *unless* the threads yield *during the increment*

In that spirit: Does this work (with 1:1 threads)?

```
1 void total(int N) {  
2     for(int i = 0; i < N; ++ i) {  
3         tally += 1;  
4         sched_yield();  
5     }  
6 }
```

Nope. It might be less likely that it is interrupted during the increment (Why? it never uses its timeslice), but it is still possible (e.g. due to

...And Adding ULTs!

What would happen if we used ULTs?

Can ULTs be pre-empted? With cooperatively-scheduled (the common case) no!

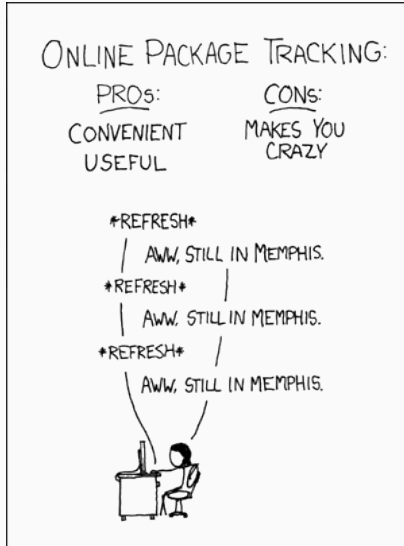
⇒ Everything will be fine, *unless* the threads yield *during the increment*

In that spirit: Does this work (with 1:1 threads)?

```
1 void total(int N) {  
2     for(int i = 0; i < N; ++ i) {  
3         tally += 1;  
4         sched_yield();  
5     }  
6 }
```

Nope. It might be less likely that it is interrupted during the increment (Why? it never uses its timeslice), but it is still possible (e.g. due to a hardware interrupt)

<https://deadlockempire.github.io>



XKCD 281 - Package Tracking

F R A G E N ?



<https://forms.gle/9CwJSKidKibubran9>

Bis nächste Woche