

Betriebssysteme

10. Tutorium - Synchronization und Deadlocks

Peter Bohner

17. Januar 2024

ITEC - Operating Systems Group

- Allgemeiner Hinweis: Was auf den ÜB drankommt ist klausurrelevant, auch wenn nicht in der VL behandelt.

- Allgemeiner Hinweis: Was auf den ÜB drankommt ist klausurrelevant, auch wenn nicht in der VL behandelt.
- T8.1a:

- Allgemeiner Hinweis: Was auf den ÜB drankommt ist klausurrelevant, auch wenn nicht in der VL behandelt.
- T8.1a: Mutual Exclusion, Progress, Bounded waiting. Was sind das?

- Allgemeiner Hinweis: Was auf den ÜB drankommt ist klausurrelevant, auch wenn nicht in der VL behandelt.
- T8.1a: Mutual Exclusion, Progress, Bounded waiting. Was sind das?
- T8.1b: Spinlock im Usermode?

- Allgemeiner Hinweis: Was auf den ÜB drankommt ist klausurrelevant, auch wenn nicht in der VL behandelt.
- T8.1a: Mutual Exclusion, Progress, Bounded waiting. Was sind das?
- T8.1b: Spinlock im Usermode? JA! Atomics (z.B. cmpxchg, fetch-and-add)

- Allgemeiner Hinweis: Was auf den ÜB drankommt ist klausurrelevant, auch wenn nicht in der VL behandelt.
- T8.1a: Mutual Exclusion, Progress, Bounded waiting. Was sind das?
- T8.1b: Spinlock im Usermode? JA! Atomics (z.B. cmpxchg, fetch-and-add)
- Interrupts ausschalten ist kein Synchronisationsprimitiv!

- Allgemeiner Hinweis: Was auf den ÜB drankommt ist klausurrelevant, auch wenn nicht in der VL behandelt.
- T8.1a: Mutual Exclusion, Progress, Bounded waiting. Was sind das?
- T8.1b: Spinlock im Usermode? JA! Atomics (z.B. cmpxchg, fetch-and-add)
- Interrupts ausschalten ist kein Synchronisationsprimitiv!
- T8.1c: Register für Spinlocks?

- Allgemeiner Hinweis: Was auf den ÜB drankommt ist klausurrelevant, auch wenn nicht in der VL behandelt.
- T8.1a: Mutual Exclusion, Progress, Bounded waiting. Was sind das?
- T8.1b: Spinlock im Usermode? JA! Atomics (z.B. cmpxchg, fetch-and-add)
- Interrupts ausschalten ist kein Synchronisationsprimitiv!
- T8.1c: Register für Spinlocks?
- P8.1: memset: alignment

- Allgemeiner Hinweis: Was auf den ÜB drankommt ist klausurrelevant, auch wenn nicht in der VL behandelt.
- T8.1a: Mutual Exclusion, Progress, Bounded waiting. Was sind das?
- T8.1b: Spinlock im Usermode? JA! Atomics (z.B. cmpxchg, fetch-and-add)
- Interrupts ausschalten ist kein Synchronisationsprimitiv!
- T8.1c: Register für Spinlocks?
- P8.1: memset: alignment
- P8.2: mq_ einmal angucken

Synchronization Primitives

There are different kinds of synchronization primitives

Which ones do you know?

Spinlock

- `lock / unlock`
- Busy-waiting and atomic instructions (e.g. compare-and-set)

There are different kinds of synchronization primitives

Which ones do you know?

Spinlock

- `lock / unlock`
- Busy-waiting and atomic instructions (e.g. compare-and-set)
- Recommended for *short* critical sections as it wastes CPU time

There are different kinds of synchronization primitives

Which ones do you know?

Spinlock

- `lock / unlock`
- Busy-waiting and atomic instructions (e.g. compare-and-set)
- Recommended for *short* critical sections as it wastes CPU time
- Preemption wastes more resources (threads can't make progress)

There are different kinds of synchronization primitives

Which ones do you know?

Spinlock

- `lock` / `unlock`
- Busy-waiting and atomic instructions (e.g. compare-and-set)
- Recommended for *short* critical sections as it wastes CPU time
- Preemption wastes more resources (threads can't make progress)
- \Rightarrow Mostly used in the kernel without interrupts

Semaphore

Semaphore

- `wait(sem)` (also called *acquire*) / `signal(sem)` (Also called *release/post*)

Semaphore

- `wait(sem)` (also called *acquire*) / `signal(sem)` (Also called *release/post*)
- Has an internal counter that is decremented (wait) or incremented (signal)

Semaphore

- `wait(sem)` (also called *acquire*) / `signal(sem)` (Also called *release/post*)
- Has an internal counter that is decremented (wait) or incremented (signal)
- Blocks if you try to decrement it below 0

Semaphore

- `wait(sem)` (also called *acquire*) / `signal(sem)` (Also called *release/post*)
- Has an internal counter that is decremented (wait) or incremented (signal)
- Blocks if you try to decrement it below 0
- Can be used to implement bounded buffers

Semaphore

- `wait(sem)` (also called *acquire*) / `signal(sem)` (Also called *release/post*)
- Has an internal counter that is decremented (wait) or incremented (signal)
- Blocks if you try to decrement it below 0
- Can be used to implement bounded buffers

Mutex

Semaphore

- `wait(sem)` (also called *acquire*) / `signal(sem)` (Also called *release/post*)
- Has an internal counter that is decremented (wait) or incremented (signal)
- Blocks if you try to decrement it below 0
- Can be used to implement bounded buffers

Mutex (Binary Semaphore)

Semaphore

- `wait(sem)` (also called *acquire*) / `signal(sem)` (Also called *release/post*)
- Has an internal counter that is decremented (wait) or incremented (signal)
- Blocks if you try to decrement it below 0
- Can be used to implement bounded buffers

Mutex (Binary Semaphore)

- `lock(m)`, `unlock(m)`

Semaphore

- `wait(sem)` (also called *acquire*) / `signal(sem)` (Also called *release/post*)
- Has an internal counter that is decremented (wait) or incremented (signal)
- Blocks if you try to decrement it below 0
- Can be used to implement bounded buffers

Mutex (Binary Semaphore)

- `lock(m)`, `unlock(m)`
- Or a Semaphore with values 0 and 1

Synchronization Primitives

Condition Variables

```
1 void consume() {
2     lock(l);
3     while(queue.size == 0) {
4         unlock(l);
5         sleep(); lock(l);
6     }
7     queue.poll(); unlock(l); signal();
8 }
9 void produce() {
10    lock(l);
11    while(queue.size == MAX_SIZE) {
12        unlock(l);
13        sleep(); lock(l);
14    }
15    queue.add(X); unlock(l); signal();
16 }
```

This code can incorrectly sleep a consumer/producer. How?

Synchronization Primitives

Condition Variables

```
1 void consume() {
2     lock(l);
3     while(queue.size == 0) {
4         unlock(l);
5         sleep(); lock(l);
6     }
7     queue.poll(); unlock(l); signal();
8 }
9 void produce() {
10    lock(l);
11    while(queue.size == MAX_SIZE) {
12        unlock(l);
13        sleep(); lock(l);
14    }
15    queue.add(X); unlock(l); signal();
16 }
```

This code can incorrectly sleep a consumer/producer. How? *Lost wakeup problem*

Synchronization Primitives

Condition Variables

```
1 void consume() {
2     lock(l);
3     while(queue.size == 0) {
4         // unlocks and sleeps atomically. Relocks when waking up
5         wait(cond_filled, l);
6     }
7     queue.poll(); signal(cond_empty); unlock(l);
8 }
9 void produce() {
10    lock(l);
11    while(queue.size == MAX_SIZE) {
12        // unlocks and sleeps atomically. Relocks when waking up
13        wait(cond_empty, l);
14    }
15    queue.add(X); signal(cond_filled); unlock(l);
16 }
```

Now no wakeup is lost :)

Deadlocks

What is that? Do you know an example?

What is that? Do you know an example?


- Several processes or activities can not make progress, as they are waiting for resources held by each other

What is that? Do you know an example?

- Several processes or activities can not make progress, as they are waiting for resources held by each other
- Examples: 4-way intersection, *Dining Philosophers*

Dining Philosophers

The Problem




Philosophers (P) want to eat, but to do that they need two Chopsticks (Chop)!

How can this deadlock?

Dining Philosophers

The Problem




Philosophers (P) want to eat, but to do that they need two Chopsticks (Chop)!

How can this deadlock?

Dining Philosophers

The Problem




Philosophers (P) want to eat, but to do that they need two Chopsticks (Chop)!

How can this deadlock?

Dining Philosophers

The Problem




Philosophers (P) want to eat, but to do that they need two Chopsticks (Chop)!

How can this deadlock?

Dining Philosophers

The Problem




Philosophers (P) want to eat, but to do that they need two Chopsticks (Chop)!

How can this deadlock?

Dining Philosophers

The Problem




Philosophers (P) want to eat, but to do that they need two Chopsticks (Chop)!

How can this deadlock?

Dining Philosophers

The Problem




Philosophers (P) want to eat, but to do that they need two Chopsticks (Chop)!

How can this deadlock?

Dining Philosophers

The Problem



Philosophers (P) want to eat, but to do that they need two Chopsticks (Chop)!

How can this deadlock?

Why did that happen? What fateful circumstances lead to this starvation?

The Four Horsemen of the Apocalypse *Coffman Conditions*

The Four Horsemen of the Apocalypse *Coffman Conditions*

Mutual Exclusion

The Four Horsemen of the Apocalypse *Coffman Conditions*

Mutual Exclusion

Resources can not be shared between processes

The Four Horsemen of the Apocalypse *Coffman Conditions*

Mutual Exclusion

Resources can not be shared between processes

Hold and wait

A process already holding resources can acquire more

The Four Horsemen of the Apocalypse *Coffman Conditions*

Mutual Exclusion

Resources can not be shared between processes

Hold and wait

A process already holding resources can acquire more

No Preemption

Resources can not be forcibly taken away from processes

The Four Horsemen of the Apocalypse *Coffman Conditions*

Mutual Exclusion

Resources can not be shared between processes

Hold and wait

A process already holding resources can acquire more

No Preemption

Resources can not be forcibly taken away from processes

Circular Wait

There exists a set of Processes P_0, P_1, \dots, P_n where P_0 is waiting for a resource held by P_1 . P_1 is waiting for a resource held by P_2 , ...and P_n is waiting for a resource held by P_1

The Four Horsemen of the Apocalypse *Coffman Conditions*

Mutual Exclusion

Resources can not be shared between processes

Hold and wait

A process already holding resources can acquire more

No Preemption

Resources can not be forcibly taken away from processes

Circular Wait

There exists a set of Processes P_0, P_1, \dots, P_n where P_0 is waiting for a resource held by P_1 . P_1 is waiting for a resource held by P_2 , ...and P_n is waiting for a resource held by P_1

Note

These conditions *are not independent!* (e.g. Circular Wait \Rightarrow Hold And Wait)

Finding a deadlock

Code

```
1  Spinlock s1,s2, s3 = FREE;
2  int counter = 0;
3  void Thread1() {
4      if(counter == 0) {
5          lock(s1);
6          counter++;
7          unlock(s1);
8      }
9      lock(s2);
10     lock(s3);
11     // update some more data
12     unlock(s3);
13     unlock(s2);
14 }
```

```
15 void Thread2() {
16     lock(s3);
17     counter++;
18     // update some data
19     if(counter == 2) {
20         lock(s2);
21         // update some more data
22         unlock(s2);
23     }
24     lock(s1);
25     // update even more data
26     unlock(s3);
27     unlock(s1);
28 }
```

Deadlock Prevention

Deadlock Prevention

Make a deadlock *impossible*!

Deadlock Prevention

Make a deadlock *impossible*! How?

Deadlock Prevention

Make a deadlock *impossible*! How? Break ≥ 1 of the four necessary conditions

Deadlock Prevention and Avoidance

Deadlock Prevention

Make a deadlock *impossible*! How? Break ≥ 1 of the four necessary conditions

Deadlock Avoidance

Deadlock Prevention and Avoidance

Deadlock Prevention

Make a deadlock *impossible*! How? Break ≥ 1 of the four necessary conditions

Deadlock Avoidance

- Deadlocks are still possible
- The resource allocator knows what resources are used by the processes
- The resource allocator denies requests that *might* lead to a deadlock

How can you negate *Mutual Exclusion*?

How can you negate *Mutual Exclusion*?

Sometimes: Spooling

How can you negate *Mutual Exclusion*?

Sometimes: Spooling

Like a Printer

- You send a job
- It is executed

⇒ Only the executor has access to the resource

Negate *Hold And Wait*

How can you negate *Mutual Exclusion*?

Sometimes: Spooling

Like a Printer

- You send a job
- It is executed

⇒ Only the executor has access to the resource

Negate *Hold And Wait*

Allocate resources atomically: All you will need or nothing

⇒ Once you have resources, you can no longer request new ones

How can you negate *No Preemption*?

How can you negate *No Preemption*?

Allow Preemption! Normally done by *multiplexing* resources (how RAM or CPU time is handled).

Not always possible

Negate *Circular Wait*

How can you negate *No Preemption*?

Allow Preemption! Normally done by *multiplexing* resources (how RAM or CPU time is handled).


Not always possible

Negate *Circular Wait*

Order resources and only allocate in the *same* order, everywhere.


Commonly used (and also in the current exercise :)

The Problem



What kind of vertices and edges are in this graph?

The Problem





What kind of vertices and edges are in this graph?

How can you detect a deadlock in there?

Resource Allocation Graphs



Some examples



Is there a deadlock in one of the graphs?

Resource Allocation Graphs

Some examples



Is there a deadlock in one of the graphs?

Yes, in the left. Right has a cycle *but no deadlock*.

Cycle \equiv Deadlock only holds if you have *one* instance of each resource

Deadlock Empire

<https://deadlockempire.github.io/>

Kernel Synchronization

How could you achieve mutual exclusion on Single-Core systems?

How could you achieve mutual exclusion on Single-Core systems?

Masking interrupts!

How could you achieve mutual exclusion on Single-Core systems?

Masking interrupts! Why?

How could you achieve mutual exclusion on Single-Core systems?

Masking interrupts! Why? Only one thread can run at a time, disabling interrupts prevents preemption (and other interrupt handlers)

How could you achieve mutual exclusion on Single-Core systems?

Masking interrupts! Why? Only one thread can run at a time, disabling interrupts prevents preemption (and other interrupt handlers)

Does this work on Multi-Core systems?

How could you achieve mutual exclusion on Single-Core systems?

Masking interrupts! Why? Only one thread can run at a time, disabling interrupts prevents preemption (and other interrupt handlers)

Does this work on Multi-Core systems?

Nope! Masking only affects the current CPU.
Additionally,

How could you achieve mutual exclusion on Single-Core systems?

Masking interrupts! Why? Only one thread can run at a time, disabling interrupts prevents preemption (and other interrupt handlers)

Does this work on Multi-Core systems?

Nope! Masking only affects the current CPU.

Additionally, another core could be in the same routine and access the same data

How would you solve that problem in the kernel?

How would you solve that problem in the kernel?

Real locks

Mutual Exclusion In The Kernel

How would you solve that problem in the kernel?

Real locks

The Big Kernel Lock™

Have a big lock *for the whole kernel*. Implications?

How would you solve that problem in the kernel?

Real locks

The Big Kernel Lock™

Have a big lock *for the whole kernel*. Implications? The kernel effectively serialises access

⇒ Can't make use of your processors if you have many syscalls

How would you solve that problem in the kernel?

Real locks

The Big Kernel Lock™

Have a big lock *for the whole kernel*. Implications? The kernel effectively serialises access

⇒ Can't make use of your processors if you have many syscalls

This removes the implementation of the big kernel lock, at last. A lot of people have worked on this in the past, so the credit for this patch should be with everyone who participated in the hunt. ([Commit message](#))

Fine Grained Locking

Fine Grained Locking

- + Only lock areas that are *relevant* for the *current* task
- ⇒ Have many small locks

Fine Grained Locking

- + Only lock areas that are *relevant* for the *current* task
- ⇒ Have many small locks
- Complex and error prone

Fine Grained Locking

- + Only lock areas that are *relevant* for the *current* task
- ⇒ Have many small locks
- Complex and error prone

Remember Spinlocks and Interrupt handlers?

Without disabling interrupts there is a problem:

Fine Grained Locking

- + Only lock areas that are *relevant* for the *current* task
- ⇒ Have many small locks
- Complex and error prone

Remember Spinlocks and Interrupt handlers?

Without disabling interrupts there is a problem: *Lockholder Preemption*

Fine Grained Locking

- + Only lock areas that are *relevant* for the *current* task
- ⇒ Have many small locks
- Complex and error prone

Remember Spinlocks and Interrupt handlers?

Without disabling interrupts there is a problem: *Lockholder Preemption*

1. Thread enters spinlock

Fine Grained Locking

- + Only lock areas that are *relevant* for the *current* task
- ⇒ Have many small locks
- Complex and error prone

Remember Spinlocks and Interrupt handlers?

Without disabling interrupts there is a problem: *Lockholder Preemption*

1. Thread enters spinlock
2. Thread gets pre-empted by an interrupt handler

Fine Grained Locking

- + Only lock areas that are *relevant* for the *current* task
- ⇒ Have many small locks
- Complex and error prone

Remember Spinlocks and Interrupt handlers?

Without disabling interrupts there is a problem: *Lockholder Preemption*

1. Thread enters spinlock
 2. Thread gets pre-empted by an interrupt handler
 3. Interrupt handler needs the same lock ⇒ Can never acquire it!
- ⇒ You might still need to disable interrupts for those

AS A PROJECT WEARS ON, STANDARDS
FOR SUCCESS SLIP LOWER AND LOWER.



FRAGEN?



<https://forms.gle/9CwJSKidKibubran9>

Bis nächste Woche

XKCD 349 - Success